

Sentiment Analysis

MohammadKazem Rajabi
Baharehsadat Khatami



Project Overview

Sentiment Analysis of Tech dataset

Objective:

- Perform sentiment analysis on the comments related to the companies like apple , lenevo and samsung .
- Classify public opinions into Positive, Neutral, and Negative categories.
- Dataset Size: ~70,000 opinions.



Project Overview

Understanding Sentiment Analysis :

- We reviewed more than 10 research papers to explore different sentiment analysis methodologies.
- Key insights from the literature:
- Strengths and weaknesses of different NLP models (e.g., LSTM, CNN, Transformers).
- Role of Attention Mechanisms in improving model performance.
- Impact of topic modeling on sentiment interpretation.



Dataset Overview

Dataset Source:

- Amazon Electronics 5-core dataset from Stanford SNAP
- Includes reviews for tech products (e.g., Apple, Samsung)

Dataset Status:

- Total reviews: ~66,768
- Brands covered: Apple, Samsung, Lenovo

Extract product reviews specifically for Apple, Samsung, and Lenovo:

- Keyword-Based Filtering:
- We searched for brand-specific keywords within review titles, product descriptions, review texts.



Dataset Construction

Keywords Used:

- Apple: apple, iphone, ipad, macbook, ipod, mac, apple pencil ...
- Samsung: samsung, galaxy, note, samsung tv, tab, fold , note , ultra...
- Lenovo: lenovo, thinkpad, ideapad, yoga,thinkcenter...

Only reviews with clear mentions of one of the target brands were kept. This ensured brand-specific sentiment analysis and topic modeling.

Sentiment classes:

- ★ 1–2 stars → Negative (19,011)
- ★ 3 stars → Neutral (17,035)
- ★ 4–5 stars → Positive (30,722)



Data Cleaning Process

A custom text-cleaning pipeline was implemented to improve data quality.

- The following steps were applied:
- Removing special characters (e.g., #, \$, %, :, ;, -).
- Eliminating numbers and URLs.
- Tokenizing text and removing stopwords.
- Standardizing text to lowercase.
- Fixing anomalies like redundant spaces.
- Ensured no missing values in key fields like review_text or star_rating.
- Removed special characters and excessive punctuation
- Excluded overly brief or generic reviews with no clear sentiment or product context.



Class Imbalance Challenge

Observed in Our imbalanced dataset:

- Positive: ~30,722
- Negative: ~19,011
- Neutral: ~17,035

Our Solution:

- Generated **synthetic neutral reviews** using large language models (LLMs)
- Prompted LLMs to produce objective, emotionless statements representing neutral sentiment e.g., “This device was released in May 2021.”
- Using Focal loss
- A stratified split



Dataset Splitting Strategy

- The dataset was divided into three sets:
- Training (80%) – Used for model learning.
- Validation (10%) – Used for hyperparameter tuning.
- Test (10%) – Used for final model evaluation.
- A stratified split ensured the results will be real .

Transformer Models for Sentiment Classification



- Transformers are deep learning models designed for sequence-to-sequence tasks.
- Introduced in the paper *"Attention Is All You Need"* by Vaswani et al. (2017).
- Uses self-attention mechanisms to process text in parallel, making it efficient and accurate.
- Commonly used in NLP tasks like translation, sentiment analysis, and text generation.

Transformers Used in This Project:

- RoBERTa
- XLNet
- DeBERTa-v3
- CardiffNLP Model

Fine-Tuning Strategy:

- Used pre-trained transformer backbones from HuggingFace
- Fine-tuned on our cleaned dataset (3-classes)
- Evaluated performance using accuracy, F1-score, and classification reports



Structural Differences:

1. RoBERTa (Robustly optimized BERT)

- Based on BERT, but removes the Next Sentence Prediction (NSP) task
- Trained on larger datasets and with longer sequences

2. XLNet (Generalized Autoregressive Model)

- Does not use masking — instead, it predicts tokens in random permutations of sequence order
- Learns bidirectional dependencies while keeping autoregressive training
- More flexible than BERT for capturing long-term context

3. DeBERTa-v3 (Decoding-enhanced BERT with disentangled attention)

- Introduces disentangled attention: separates content (word meaning) and position (word order) in attention computation
- Uses relative position encoding, unlike BERT's absolute positions

4. CardiffNLP (Twitter-RoBERTa)

- Same architecture as RoBERTa, but pretrained on Twitter data
- Specialized for short, informal, and noisy text



Tokenization and Dataset Conversion

Understanding Tokenization

Tokenization is the process of converting text into numerical representations that a model can understand.

Transformers require subword tokenization to handle out-of-vocabulary words and improve efficiency.

We used Hugging Face's pre-trained tokenizers to tokenize our dataset.

Tokenizers Used

DistilBERT Tokenizer (distilbert-base-uncased)

RoBERTa Tokenizer (roberta-base, roberta-large)

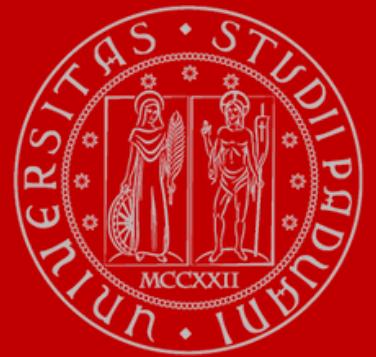
DeBERTa Tokenizer (microsoft/deberta-v3-small, microsoft/deberta-large)

How Tokenization Works

Token Splitting: Converts text into smaller units (words/subwords).

Encoding: Converts tokens into numerical IDs based on a pre-defined vocabulary.

Tokenization and Dataset Conversion



Model	Tokenizer Type	Hugging Face Class
RoBERTa	Byte-Pair Encoding (BPE)	RobertaTokenizer
XLNet	SentencePiece (Unigram LM)	XLNetTokenizer
DeBERTa-v3	WordPiece	DebertaV2Tokenizer
CardiffNLP	BPE (Tweet-optimized)	AutoTokenizer.from_pretrained

RoBERTa Base & Large



Robustly optimized version of BERT.

- Pretrained Models: roberta-base, roberta-large
- Differences from BERT:
- More training data.
- No Next Sentence Prediction (NSP).
- Uses dynamic masking.

Fine-Tuning:

- 20 epochs, batch size of 16.
- Learning rate: 2e-5 (base), 1e-5 (large).
- use early stopping

DeBERTa Small & Large



- Uses disentangled attention to improve contextual understanding.
- Pretrained Models: deberta-v3-base, deberta-large

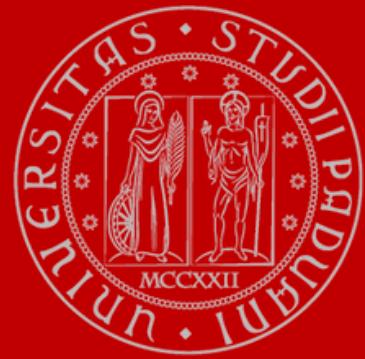
Key Features:

- Enhanced self-attention mechanism.
- Better handling of word relationships.

Fine-Tuning:

- DeBERTa-Small: 20 epochs, batch size of 16.
- DeBERTa-Large: 20 epochs, batch size of 16.
- Learning rate: 1e-5.
- use early stopping

Comparison of Model Performance(tech dataset)



Model	Accuracy	Precision	Recall	F1-score
CardiffNLP	0.7598	0.7524	0.7598	0.7546
DeBERTa-v3-large	0.7526	0.7539	0.7526	0.7532
DeBERTa-v3-base	0.7228	0.7572	0.7228	0.7332
RoBERTa	0.726	0.7204	0.726	0.7223
XLNet	0.7088	0.7371	0.7088	0.718

Conclusion



- Transformer models are highly effective for sentiment analysis.
- CardiffNLP was the most successful model in this project.

Sentiment Analysis Using Deep Learning



Models Implemented:

- LSTM
- BiLSTM
- BiLSTM + GloVe
- BiLSTM + GloVe + Attention
- BiLSTM + GloVe + Attention + Focal Loss
- GRU + GloVe
- Conv1D + BiLSTM



Training Pipeline

Training Setup:

- Train/val/test split with stratification
- Class weighting for imbalance
- Focal Loss used in final model

Optimization:

- Optimizer: Adam
- EarlyStopping on validation loss (patience = 5)

Metrics:

- Accuracy
- F1-Score (macro)
- Confusion Matrix



Model Architectures

LSTM for Sentiment Analysis

LSTM (Long Short-Term Memory) is a type of recurrent neural network that learns long-range dependencies in sequential data. It's ideal for modeling the flow and context of review text.

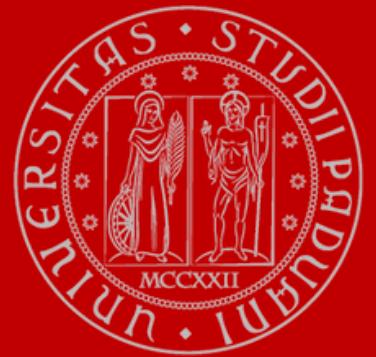
Our Model Setup:

- 2 stacked LSTM layers (256 and 128 units)
- Dropout applied to reduce overfitting
- Final dense layer with softmax for 3-class classification

Fine-Tuning Details:

- Input: Tokenized and padded review sequences
- Loss Function: Categorical Cross-Entropy
- Optimizer: Adam
- Training: 30 epochs with EarlyStopping
- Evaluation: Accuracy and macro F1-score

Model Architectures



How BiLSTM Works (Without Attention)

1. Input Sequence
battery | life | is | amazing | but | screen | is | weak

2. Embedding Layer
→ Convert tokens into 300D word vectors

3. BiLSTM Layer 1 (return_sequences=True)
→ Outputs a vector for each word (8×256)

4. BiLSTM Layer 2 (return_sequences=False)
→ Reads full sequence and returns last hidden state

5. Sentence Vector
→ Last hidden state = 128D forward + 128D backward → 256D

6. Dense + Softmax
→ Predict sentiment class based on sentence vector



Model Architectures

How Attention Works in BiLSTM + Attention

1. Input Sequence

battery | life | is | amazing | but | screen | is | weak

2. BiLSTM Output

→ A vector for each word, capturing context (8×256)

3. Attention Scores

→ Compute similarity between all word pairs ($Q \cdot K^t$)

4. Attention Weights

→ Apply softmax to highlight key words (e.g., 'amazing', 'weak')

5. Context Vector

→ Weighted sum of word vectors (focus on important words)

6. Concatenation

→ Combine [Original BiLSTM Output] + [Attention Output] → richer features

7. Final BiLSTM

→ Compress into single sentence vector (128D) for classification

Model Architectures



LSTM (Long Short-Term Memory) is a type of recurrent neural network that learns long-range dependencies in sequential data. It's ideal for modeling the flow and context of review text.

- BiLSTM:
- Adds backward processing to better capture full context in both directions.
- BiLSTM + GloVe:
- Integrates 300D pre-trained GloVe embeddings to inject external word knowledge.
- BiLSTM + GloVe + Attention:
- Adds an attention mechanism to focus on key informative words in each review.
- BiLSTM + GloVe + Attention + Focal Loss:
- Final model variant that addresses class imbalance by prioritizing harder examples.

Fine-Tuning Shared Setup:

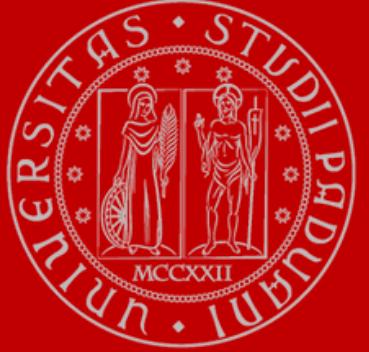
- Tokenized and padded sequences
- Adam optimizer + EarlyStopping
- Metrics: Accuracy and macro F1-score



Using Pre-Trained Embeddings

- GloVe (Global Vectors for Word Representation):
- Pre-trained word embeddings that capture word relationships and semantic meanings.
- Allowed the model to leverage linguistic knowledge from large external corpora.
- Why Pre-Trained Embeddings?
- Helps the model understand words even if they were infrequent in the training dataset.
- Improves generalization and reduces the risk of overfitting.
- Implementation:
- Integrated 300-dimensional GloVe embeddings in the embedding layer of the model.

Enhancing with GloVe & Focal Loss



GloVe Embeddings

- Pre-trained word vectors that capture semantic meaning.
- We integrated 300D GloVe embeddings into our BiLSTM models to improve word understanding.

Focal Loss

- A custom loss function designed to handle class imbalance by focusing more on hard-to-classify samples.
- Used in our final model to boost neutral class performance.

Result:

- Adding GloVe and Focal Loss led to noticeable improvements in both accuracy and macro F1-score, especially on underrepresented classes.

Model Architectures



BiLSTM with GloVe Embeddings

1. Input Review

→ Raw text: "Battery lasts all day, but display is dull"

2. Tokenization + Padding

→ Convert to token IDs: [12, 45, 7, ..., 0, 0]

3. GloVe Embedding Layer

→ Map each token to a 300D pre-trained vector (fixed meaning)

4. BiLSTM Layer 1 (return_sequences=True)

→ Context-aware word vectors (sequence of hidden states)

5. BiLSTM Layer 2 (return_sequences=False)

→ Final hidden state becomes sentence vector (128D × 2)

6. Dense + Softmax

→ Predict sentiment: Negative, Neutral, Positive



Models Performance

Model	Accuracy	F1-score
GRU + GloVe	0.7084	0.7099
BiLSTM + Attention + GloVe	0.72	0.69
BiLSTM + Attention	0.703	0.68
BiLSTM + Attention + GloVe + FL	0.7	0.65
Vanilla BiLSTM	0.6674	0.6596
GRU	0.6348	0.6408

Final Model Performance

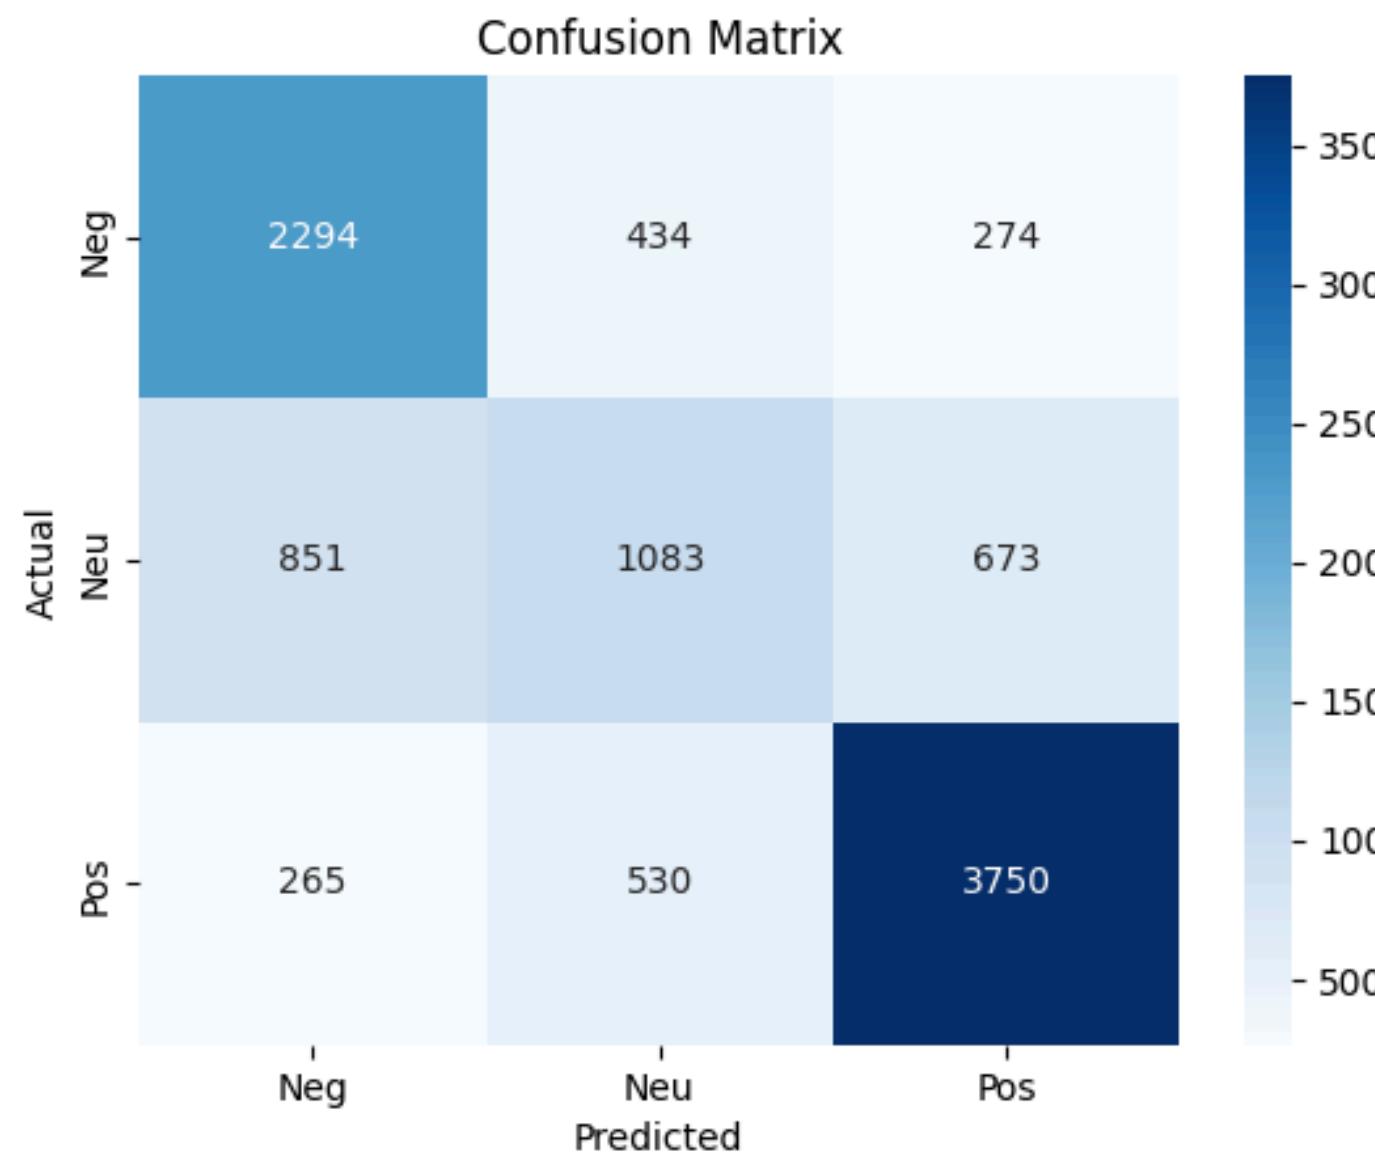


Sentiment	Precision	Recall	F1-score
Negative	0.67	0.76	0.72
Neutral	0.53	0.42	0.47
Positive	0.8	0.83	0.81



Model Performance

- The model performs best on the positive class.
- Performance on neutral sentiment is lowest, highlighting class imbalance difficulty.
- Focal Loss helped recover some performance for the minority class.



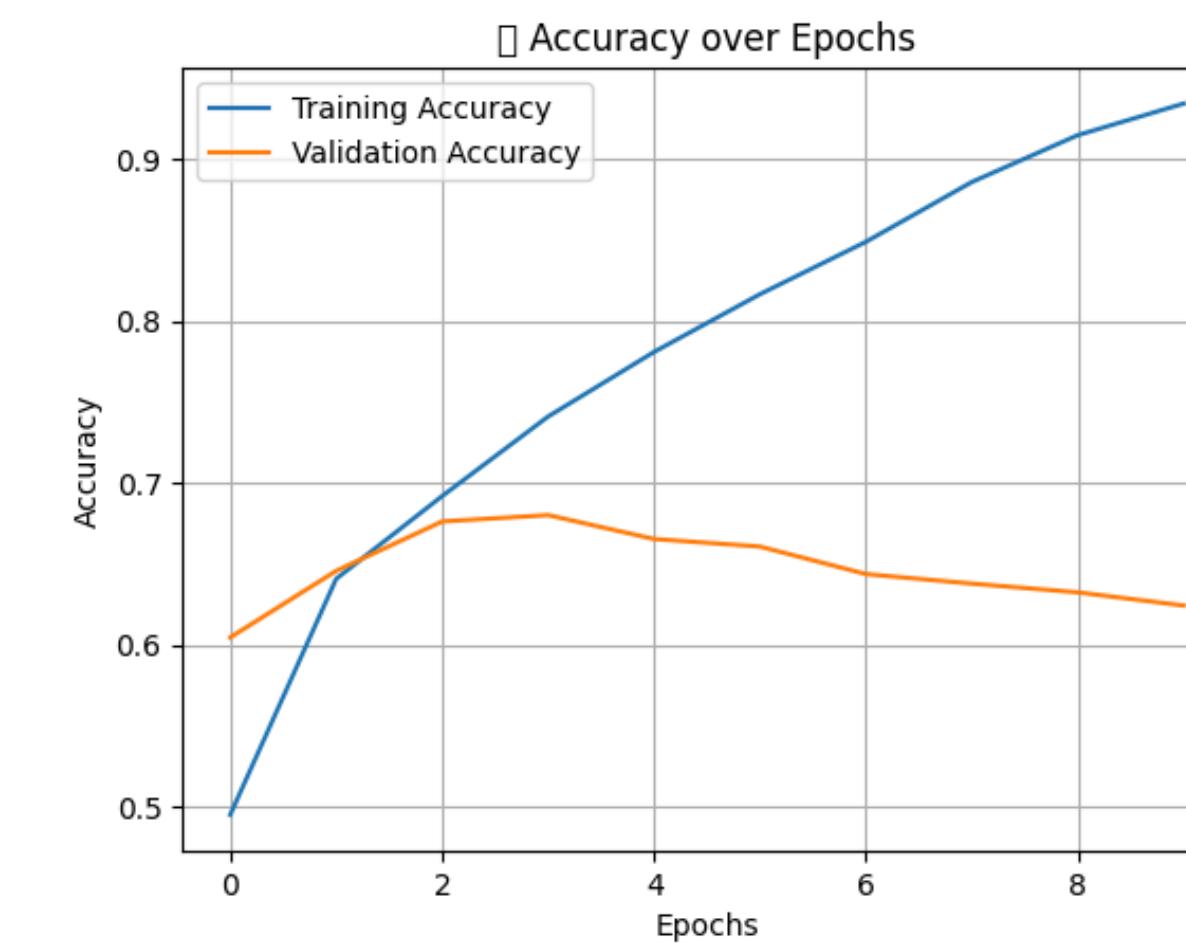
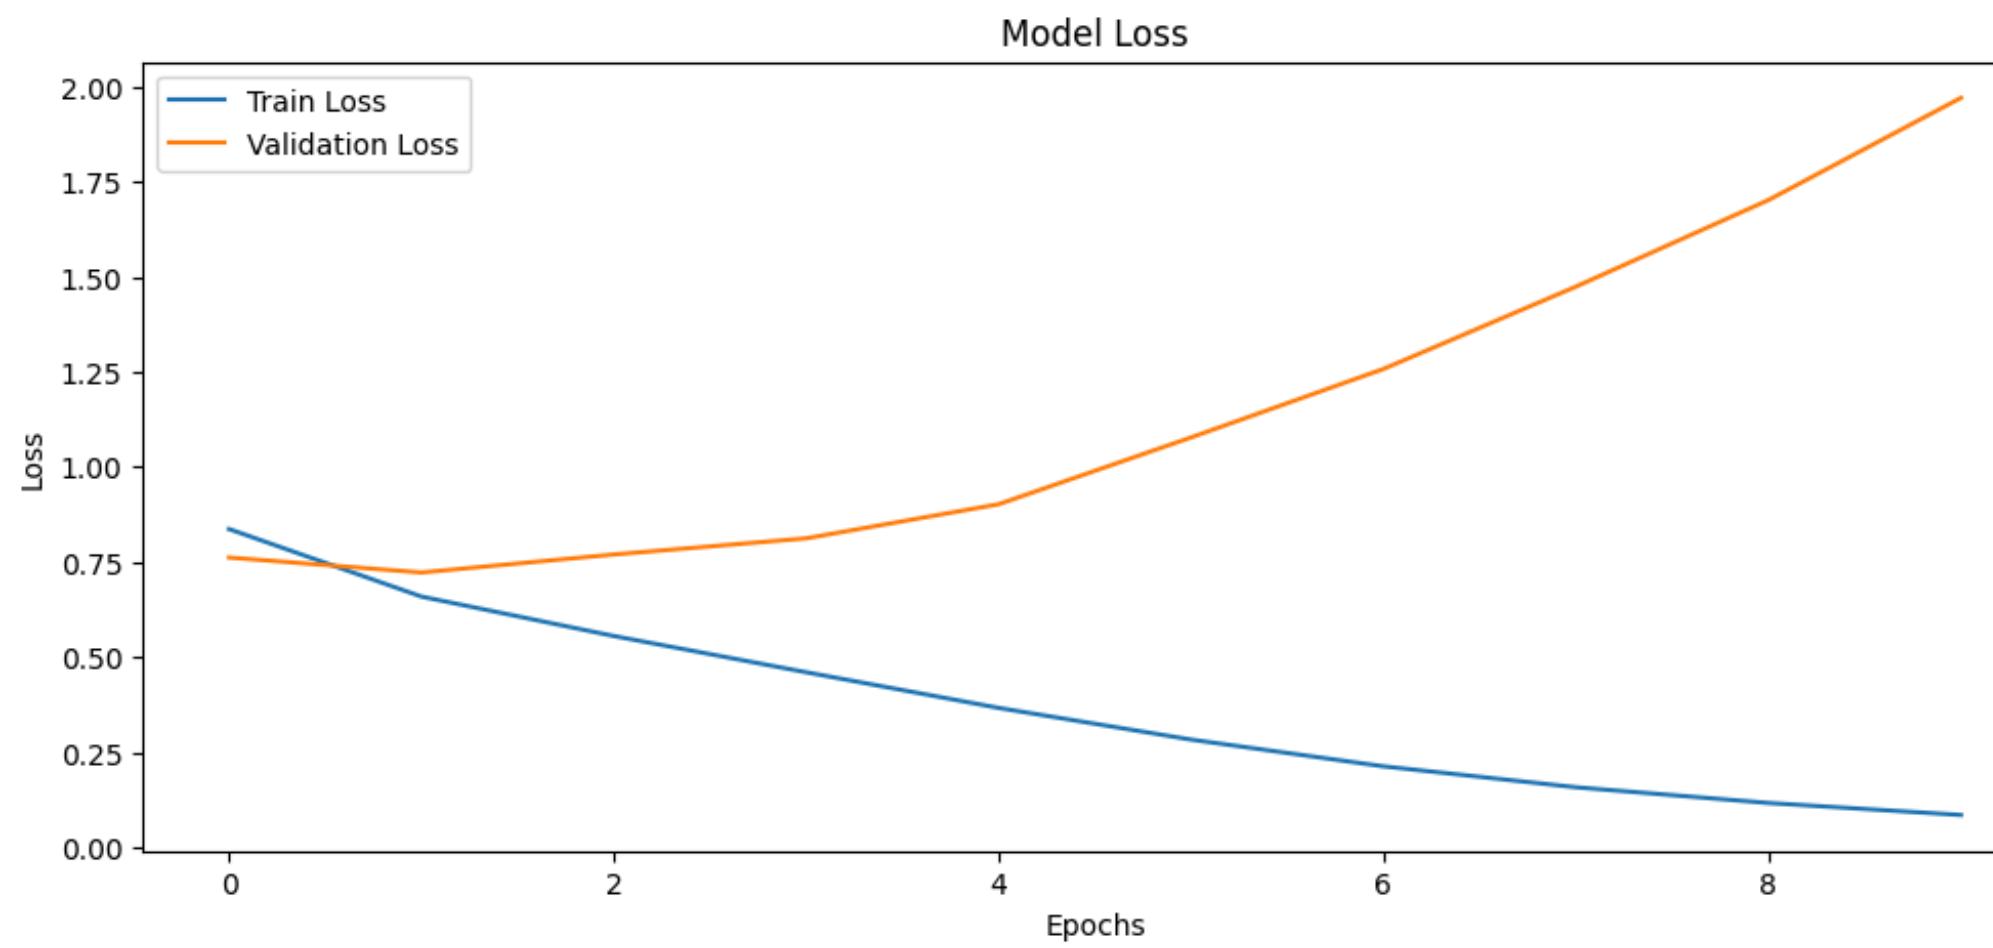
- The Galaxy S24 is fast but heats up.
→ Probabilities: [0.08534862 0.49317715 0.4214742]
→ Predicted Label: 0
- iPhone 15 Pro Max is incredible!
→ Probabilities: [0.05532453 0.08167563 0.86299986]
→ Predicted Label: 1
- Not good, not bad, just average.
→ Probabilities: [0.4101152 0.52455896 0.06532583]
→ Predicted Label: 0

Conclusion

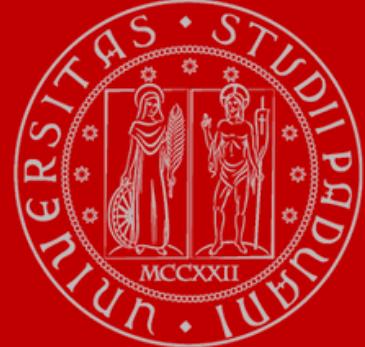


We observed that even basic models like GRU and vanilla LSTM delivered reasonable performance, especially when combined with GloVe embeddings. However, with the addition of attention mechanisms, focal loss, and transformer architectures.

The shallow GRU and LSTM models reached around 50% accuracy and F1-score, which is not sufficient for reliable sentiment classification – especially given the poor performance on the neutral class.

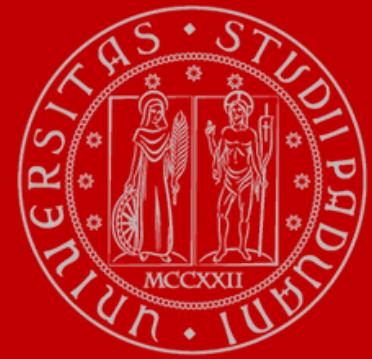


FastText: An Efficient and Accurate Text Classification Model



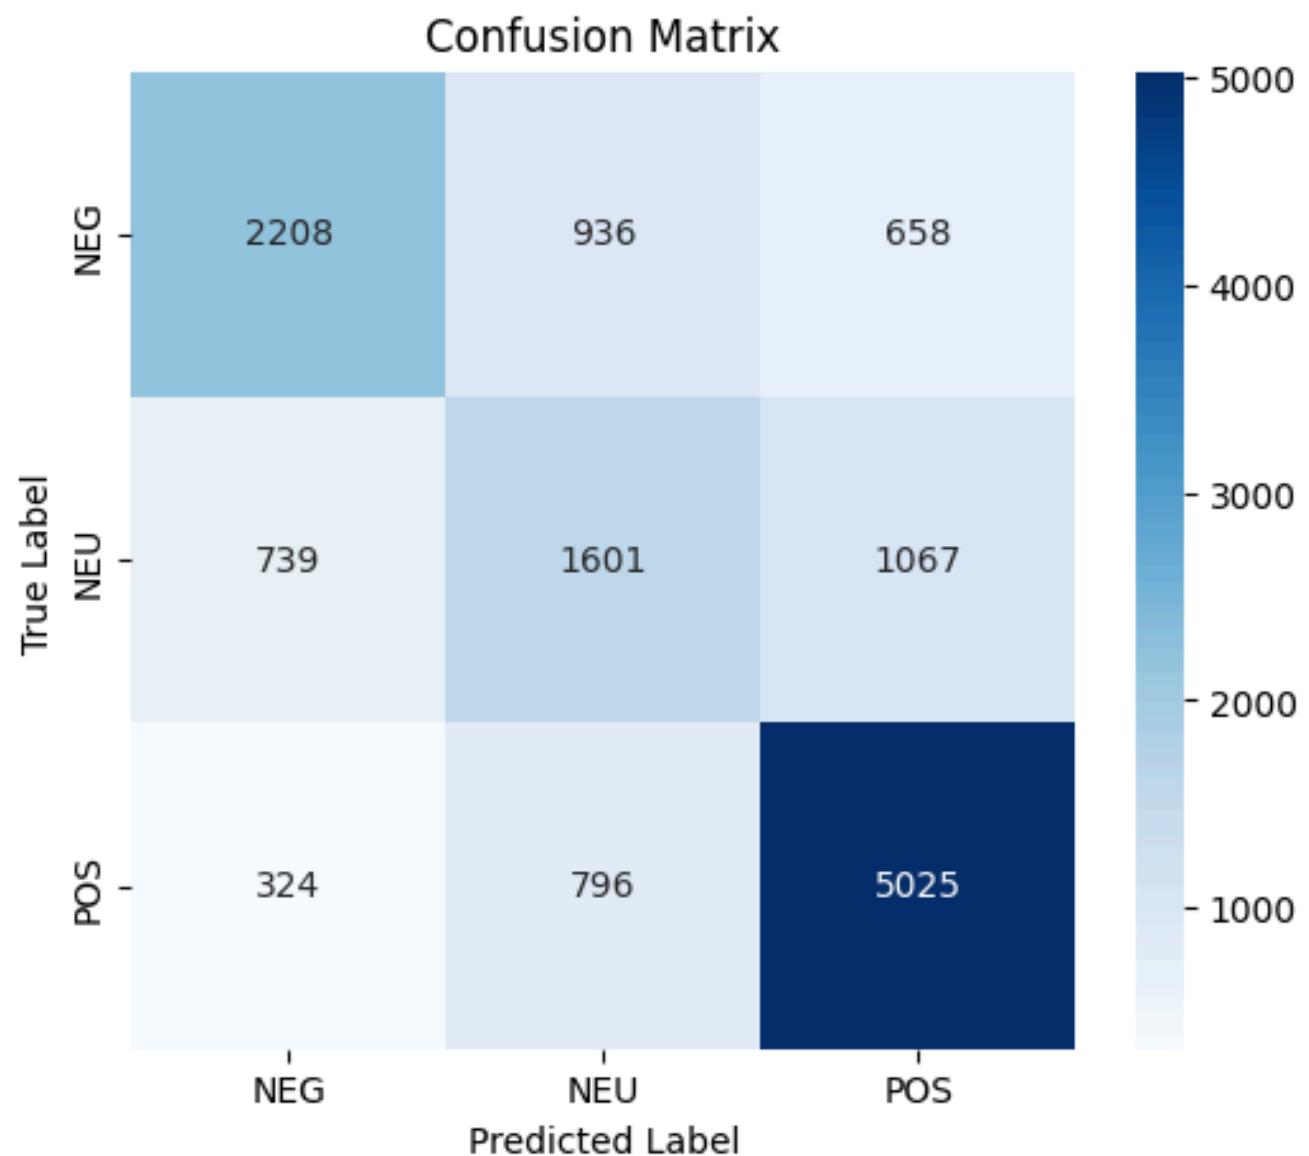
Introduction to FastText

- FastText is an open-source, efficient text classification and representation learning library developed by Facebook AI Research (FAIR). It is widely used for text classification, word embeddings, and sentence representations. FastText provides:
- Fast training and inference
- Support for n-gram features
- Ability to handle rare and out-of-vocabulary words effectively
- Suitability for multi-label classification tasks
- In our project, we leverage FastText for sentiment classification.



FastText on Our Dataset(tech dataset)

Averaged F1_score:62%



Class	Precision	Recall	F1-score
Negative	0.67	0.58	0.62
Neutral	0.48	0.46	0.47
Positive	0.74	0.8	0.76

Conclusion



- FastText is an efficient, scalable, and accurate model for text classification.
- Despite its simplicity, FastText remains a powerful choice for real-time applications where speed is crucial.

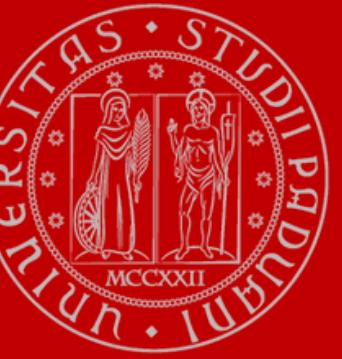
Topic Modeling



- Topic modeling is a suite of algorithms used to uncover hidden thematic structures in a collection of texts. It assigns topics to documents, enabling better organization and understanding of unstructured data
- Here the goal is to compare different topic models with different algorithms and make a conclusion
- At last we do a Semantic Distribution Among Topics for the best topic model

Models Used:

- LDA (Latent Dirichlet Allocation)
Probabilistic generative model capturing word-topic distributions.
- LSA (Latent Semantic Analysis)
Matrix decomposition using TF-IDF + SVD to discover latent patterns.
- NMF (Non-negative Matrix Factorization)
Parts-based representation of documents with non-negative constraints.
- BERTopic
Transformer-based model combining BERT embeddings, UMAP, HDBSCAN, and class-based TF-IDF for high-quality topics.



LSA uses linear algebra to uncover latent structures in the term-document matrix, assuming similar words occur in similar contexts.

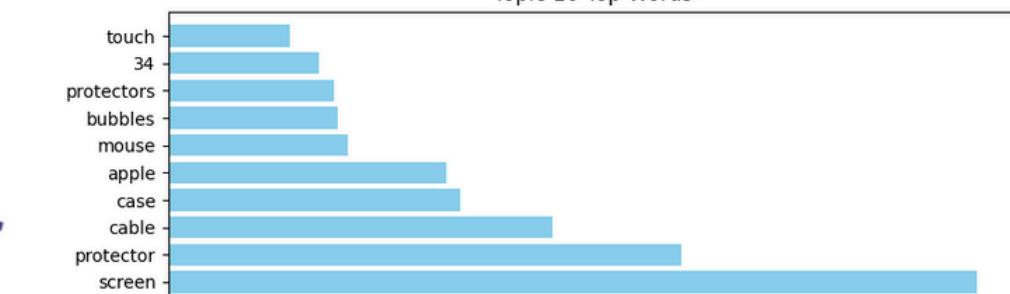
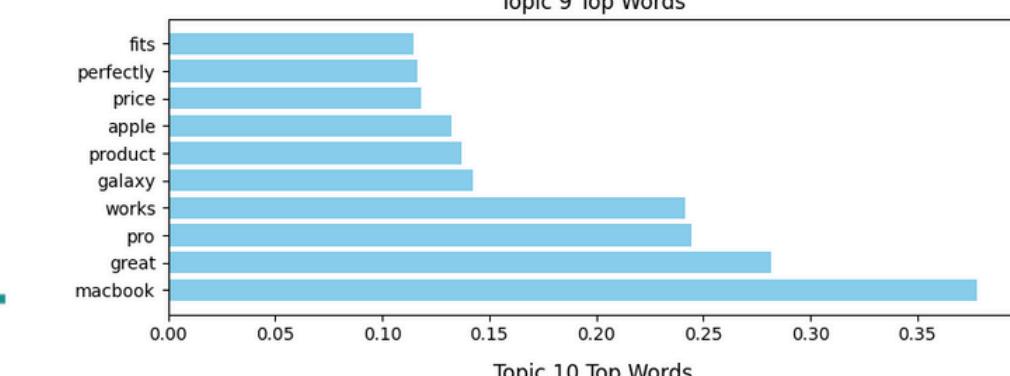
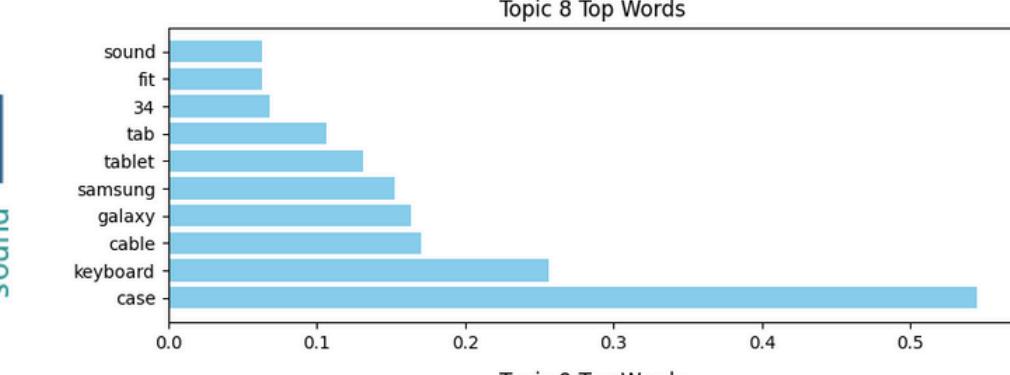
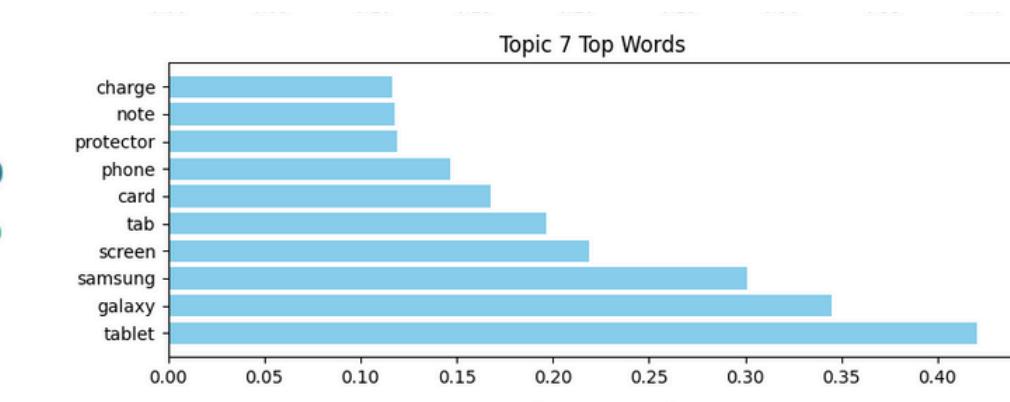
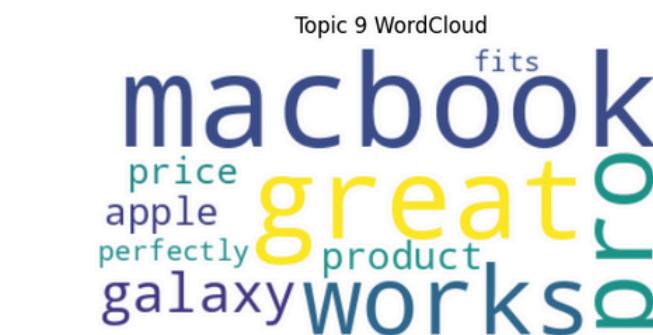
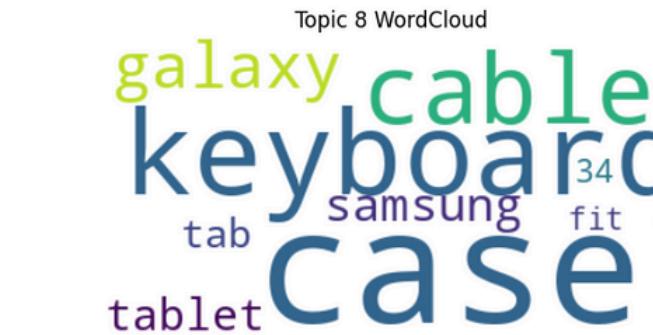
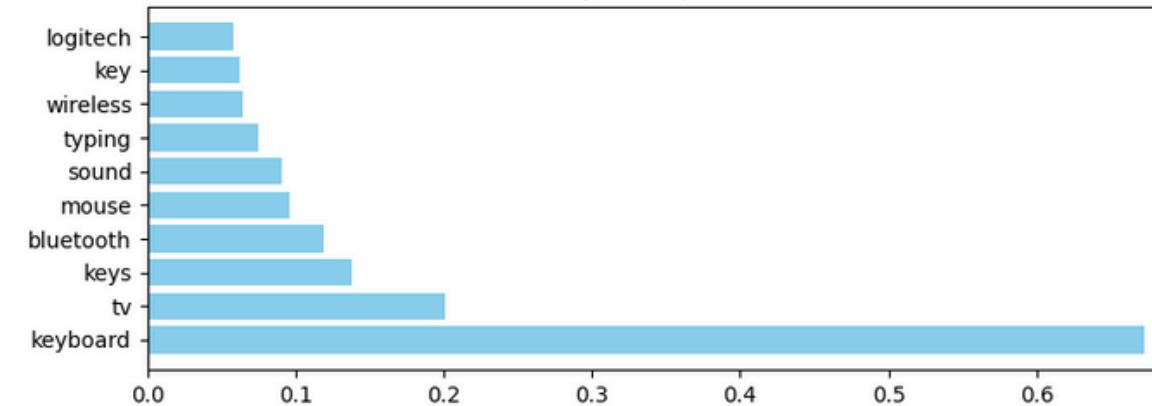
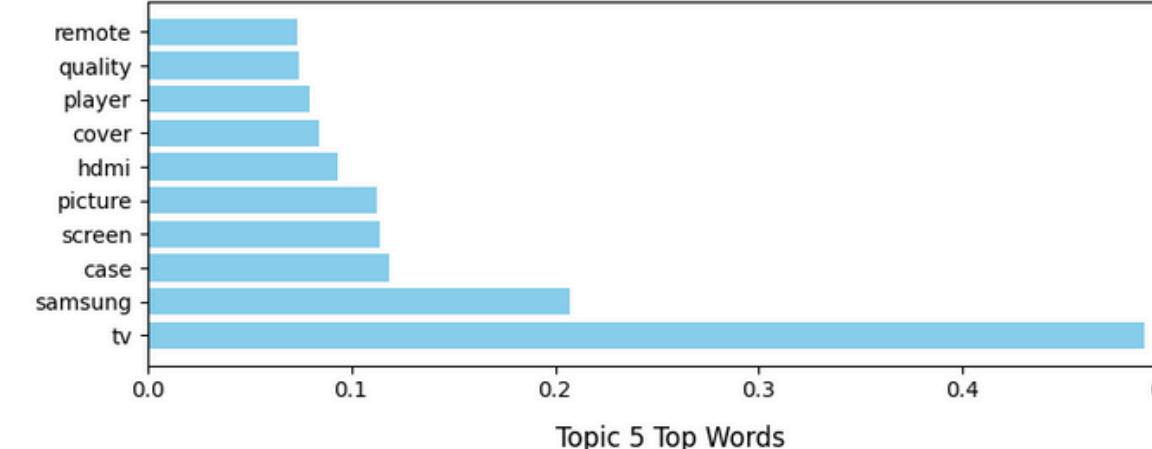
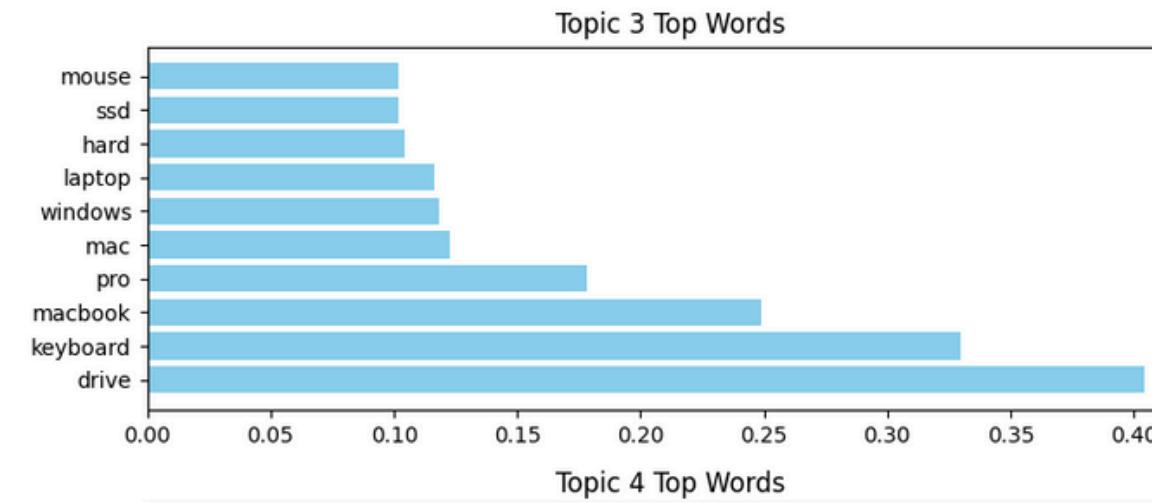
Advantages:

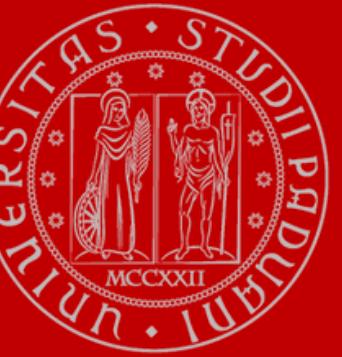
- Captures synonymy and polysemy.
- Simple and computationally efficient.

Limitations:

- Assumes linear relationships.
- Sensitive to noise in the data.

Results





NMF decomposes a non-negative term-document matrix into two non-negative factors.

Advantages:

- Topics are interpretable due to non-negativity.
- Effective for sparse data.

Limitations:

- Sensitive to initialization.
- May converge to local minima

Results



WordCloud for Topic 6

keyboard
laptop
windows
use
keys
mouse
bluetooth
key
wireless
logitech
iphone
port
device
charger
adapter

WordCloud for Topic 1

stand
mini
ipad
case
protection
like
cover
leather
original
battery
hours
replacement
fits

WordCloud for Topic 4

charge
phone
usb
charger
power
charging
volume
speaker
iphone
great
good

WordCloud for Topic 3

laptop
lenovo
hours
battery
fit
life
thinkpad
replacement
fits

WordCloud for Topic 8

headphones
sound
quality
volume
speaker
music
great
good

WordCloud for Topic 2

drives
external
ssd
usb
laptop
hard
drive
sata
windows



BERTopic leverages transformer-based embeddings and clustering techniques to extract dynamic and contextual topics.

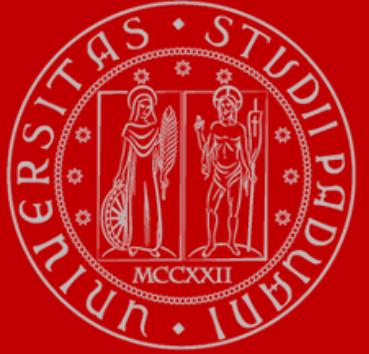
Advantages:

- Captures contextual semantics.
- Works well with short texts.

Limitations:

- Requires pre-trained embeddings.
- Computationally expensive.

Results



WordCloud for Topic 5

wifi. network
wireless
linksys signal
router modem
netgear connection support

WordCloud for Topic 3

lenovo processor
laptops ram
hp laptop
computer screen
machine windows

WordCloud for Topic 0

stand cover is
. this i pad
for and the case

WordCloud for Topic 4

video zoom images
cameras camera images
camera lens photos pictures
canon

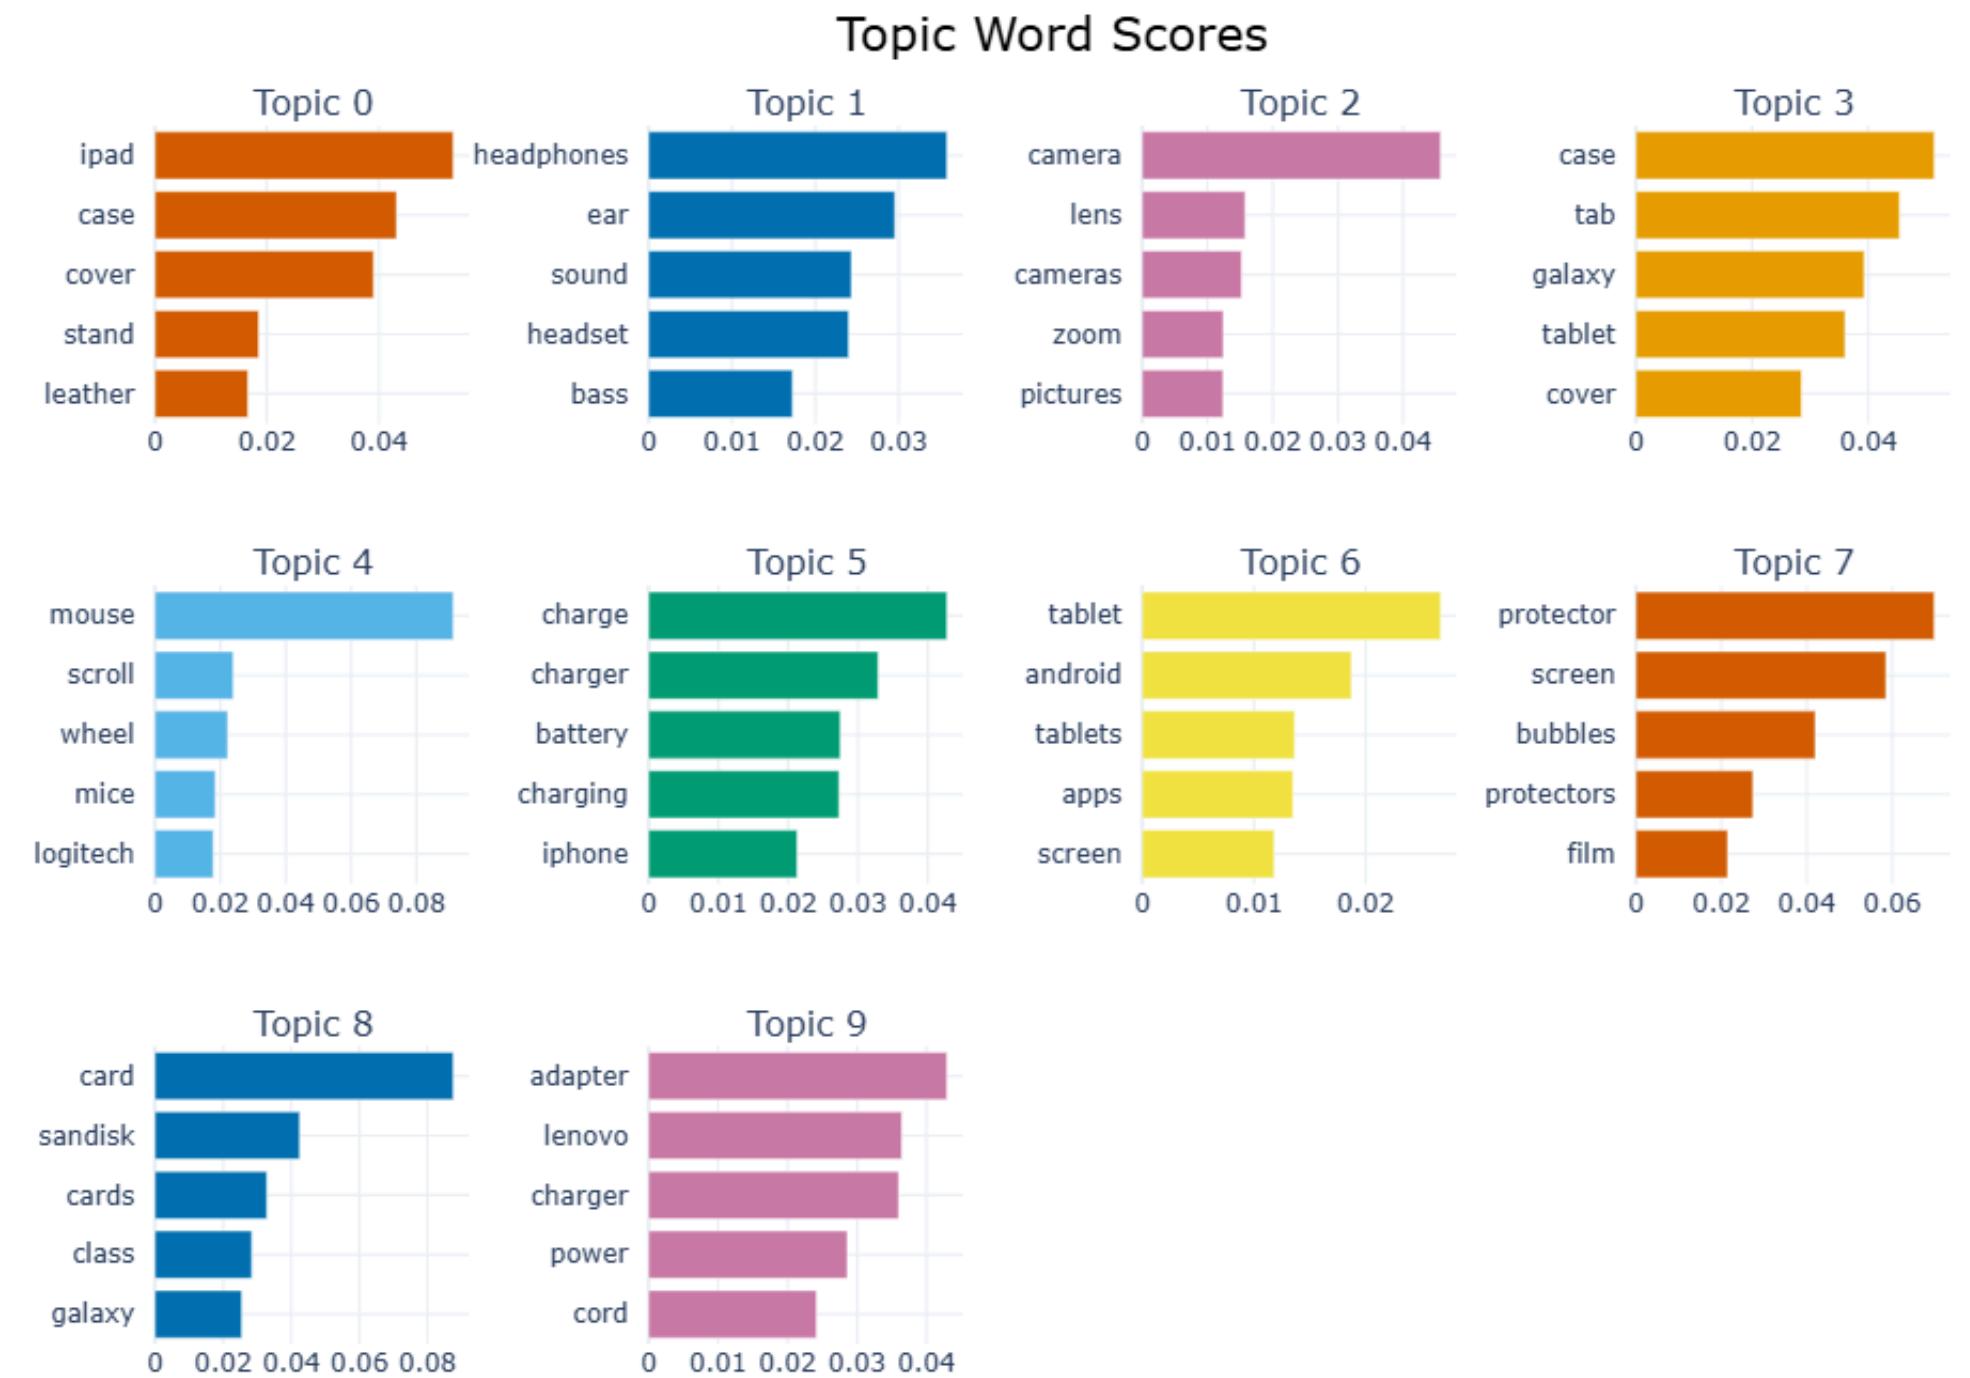
WordCloud for Topic 8

hand mice buttons
mouse button microsoft
wheel scroll logitech
bluetooth

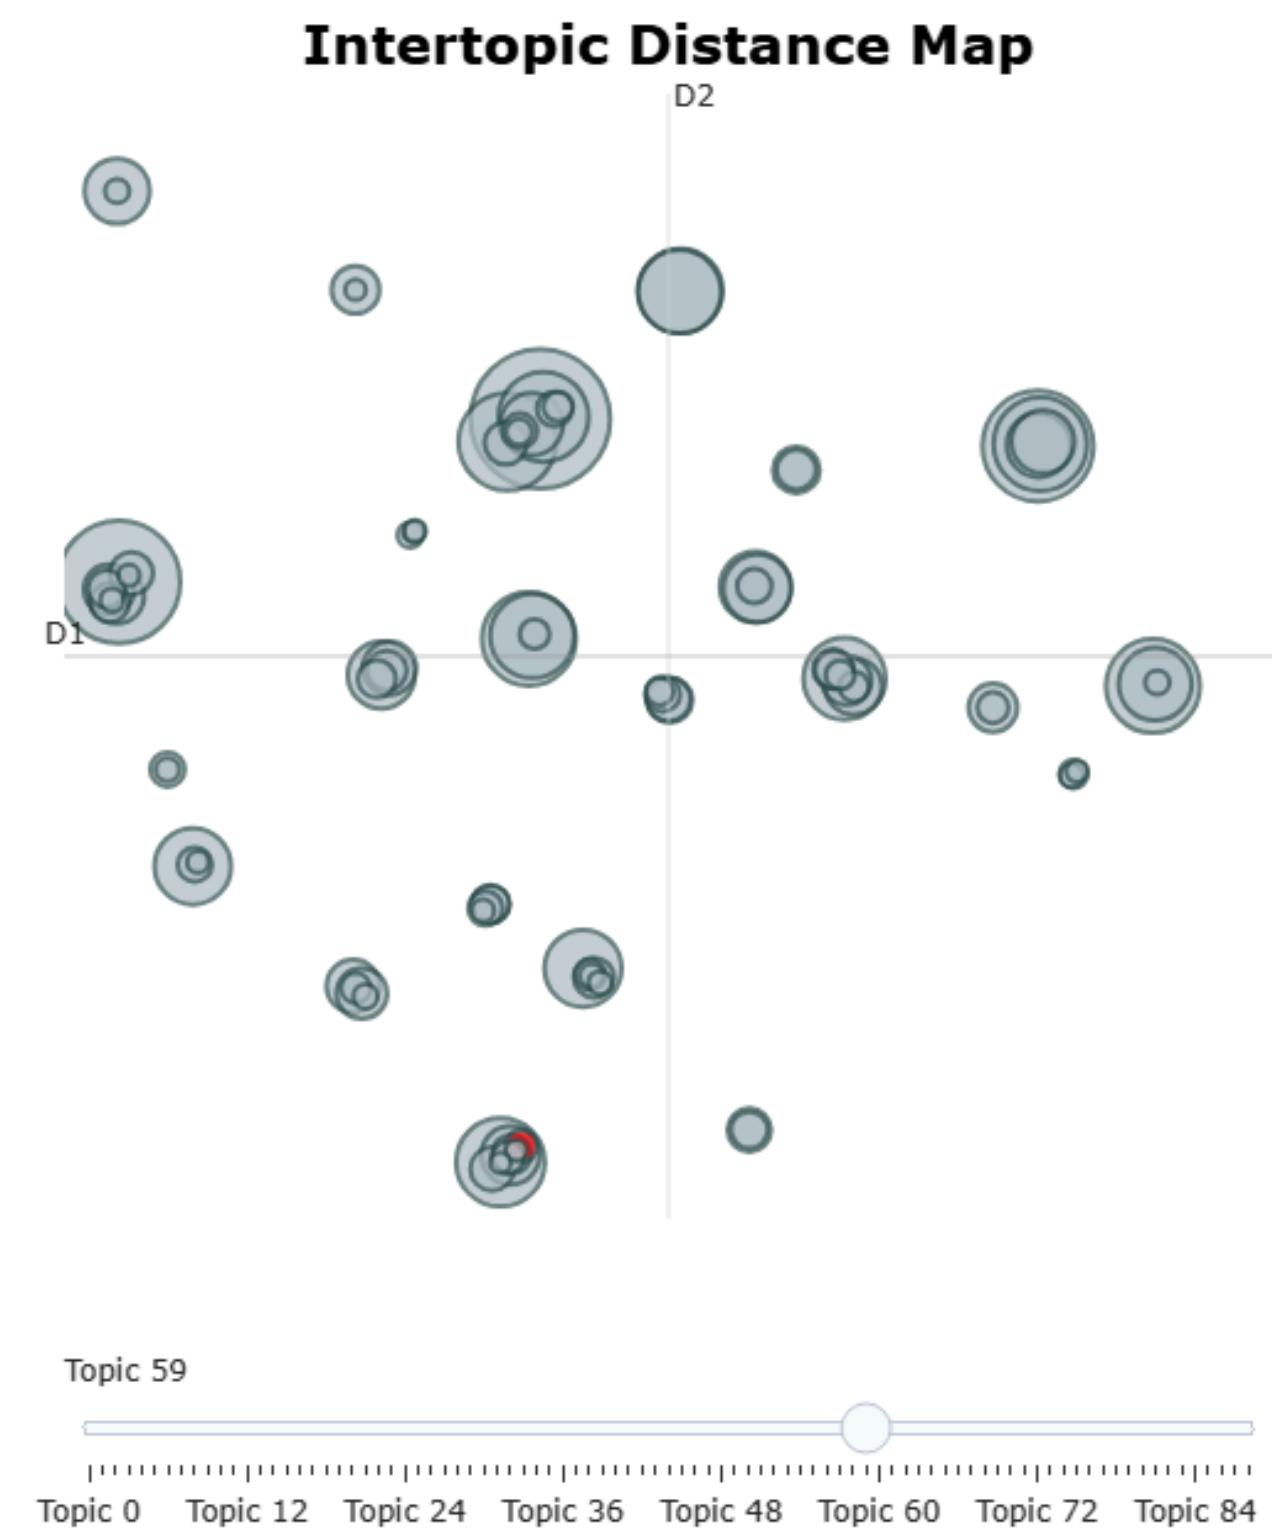
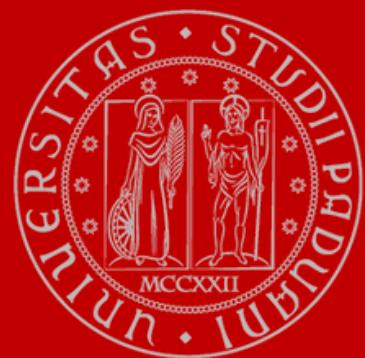
WordCloud for Topic 13

dvi display cable
displayport connect laptop
monitor lenovo adapter

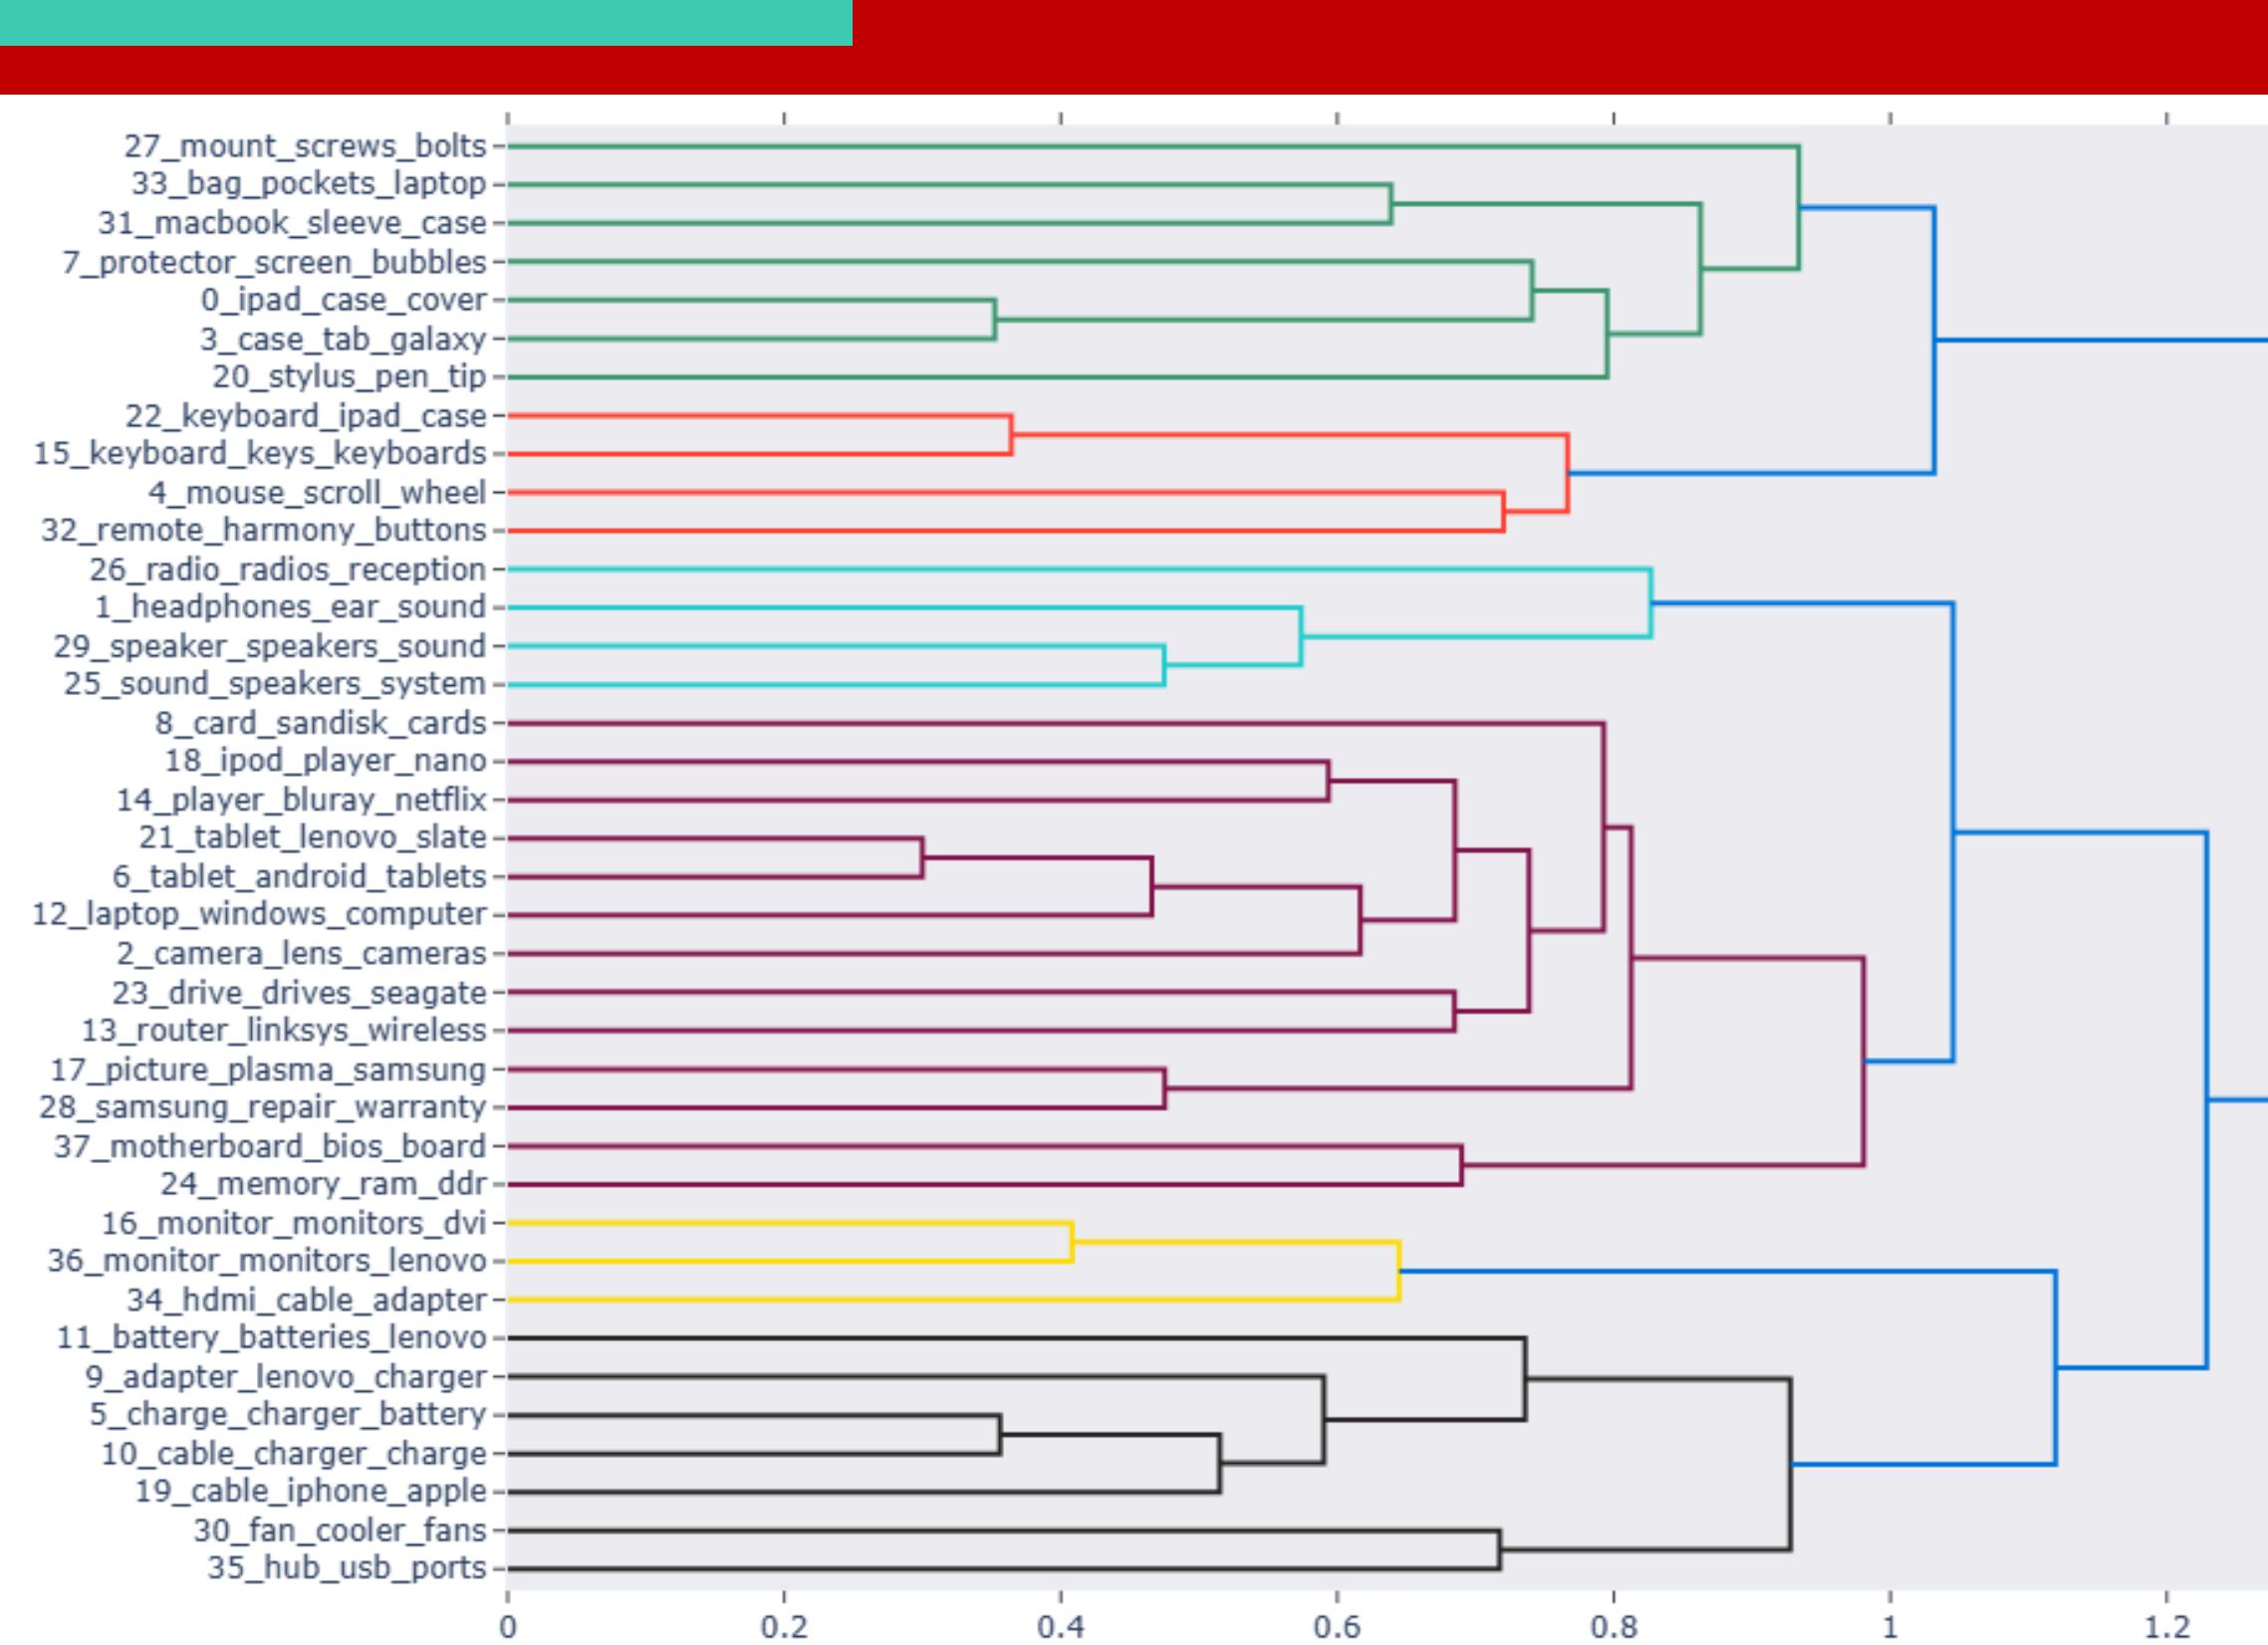
Results

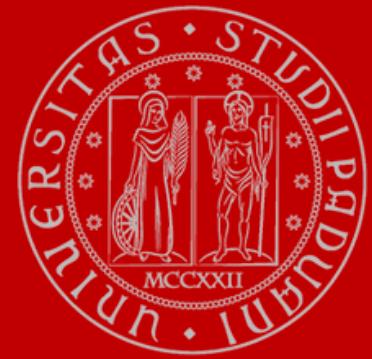


Results

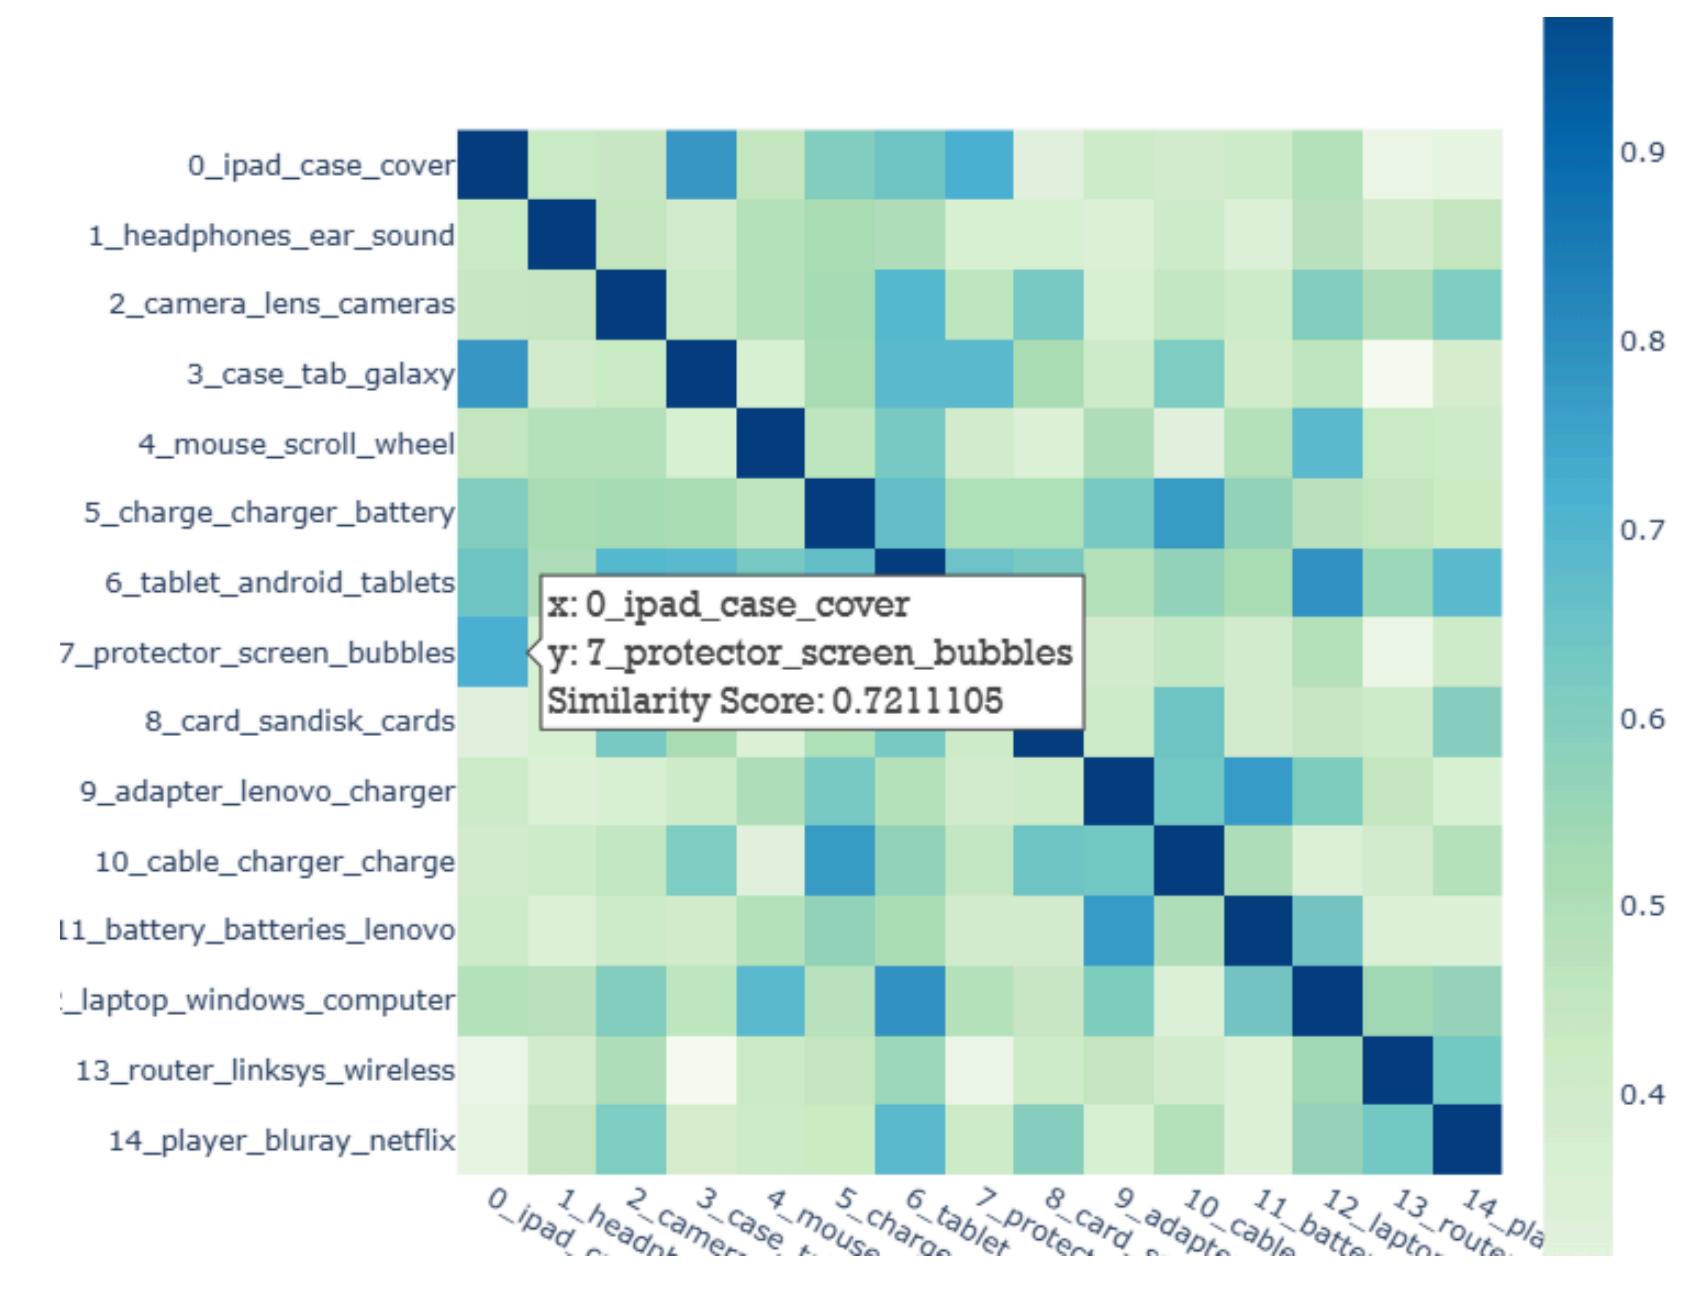
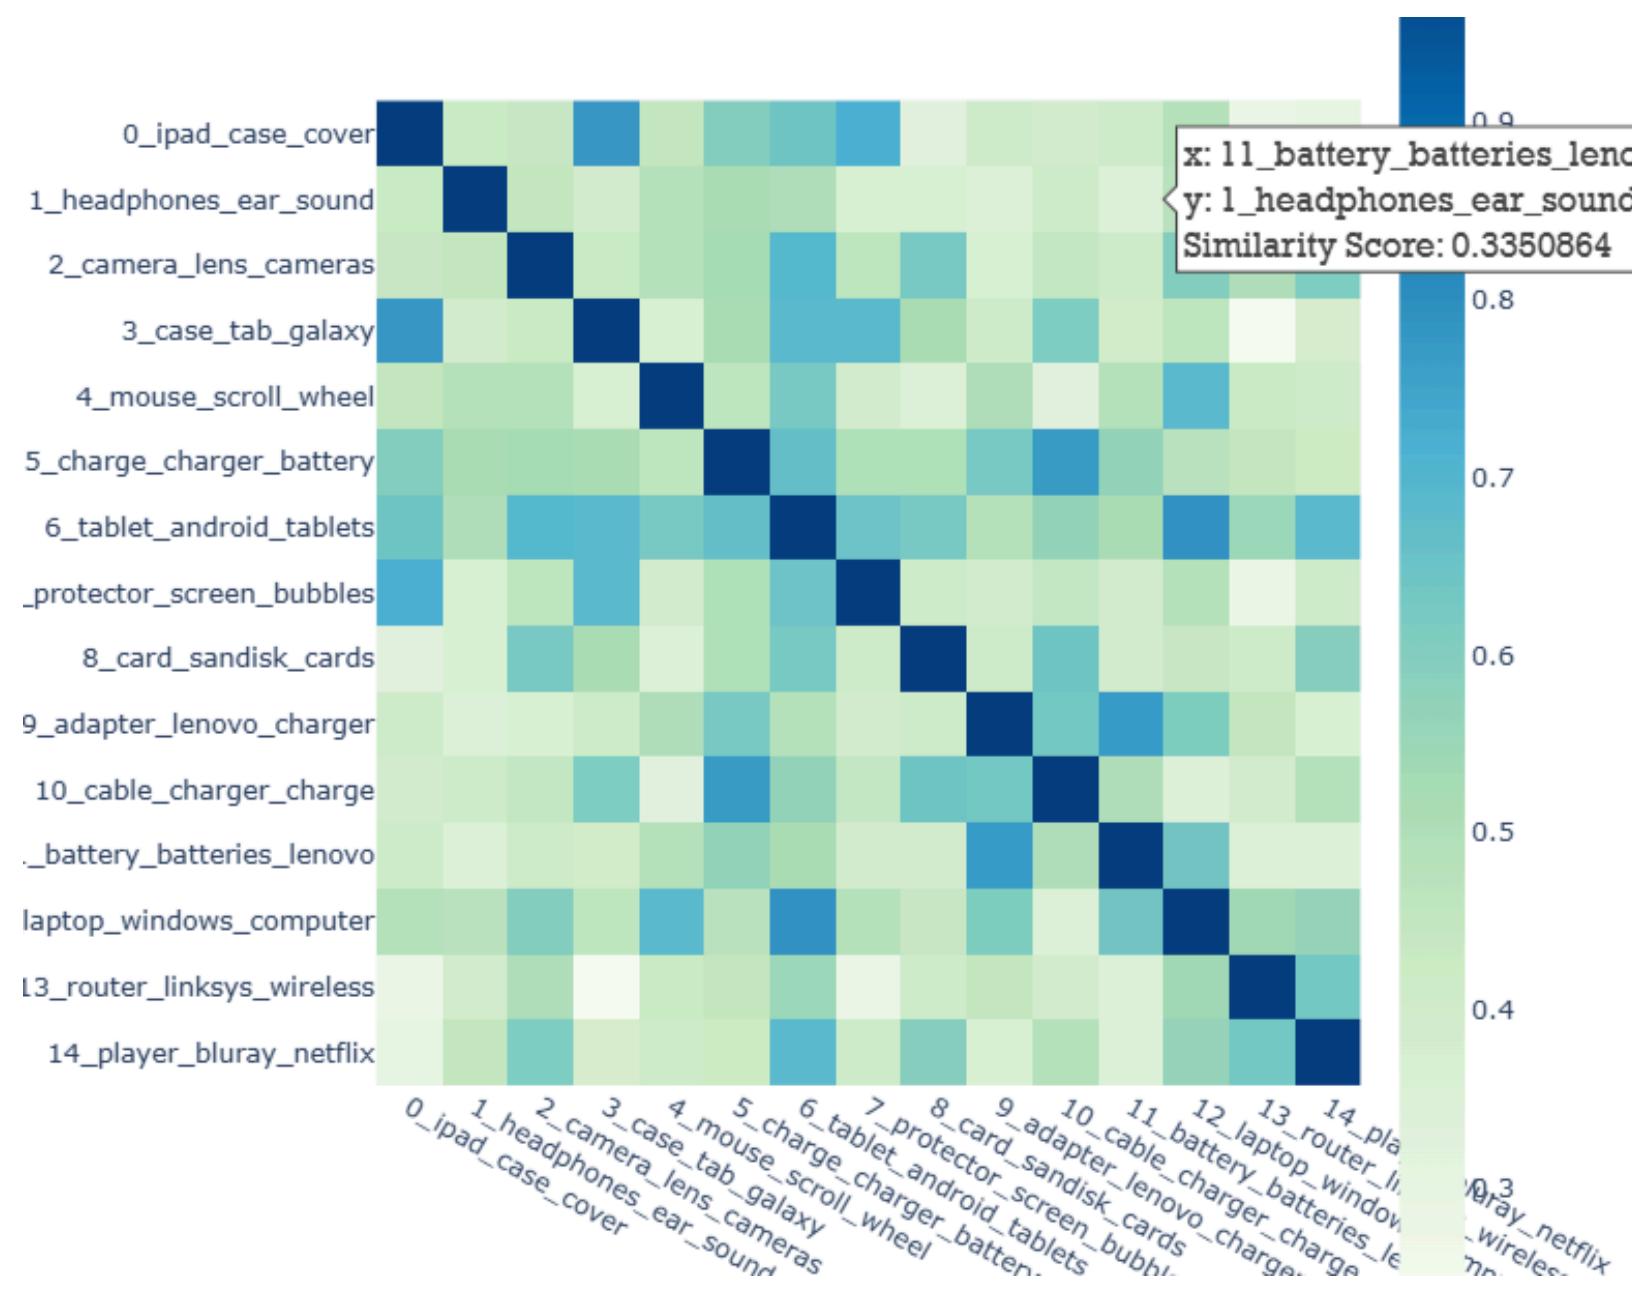


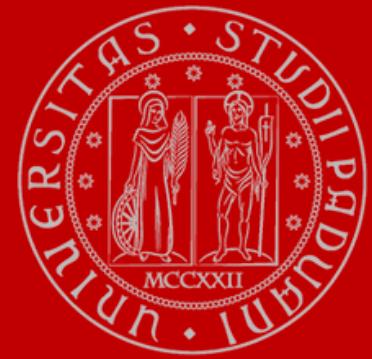
Results



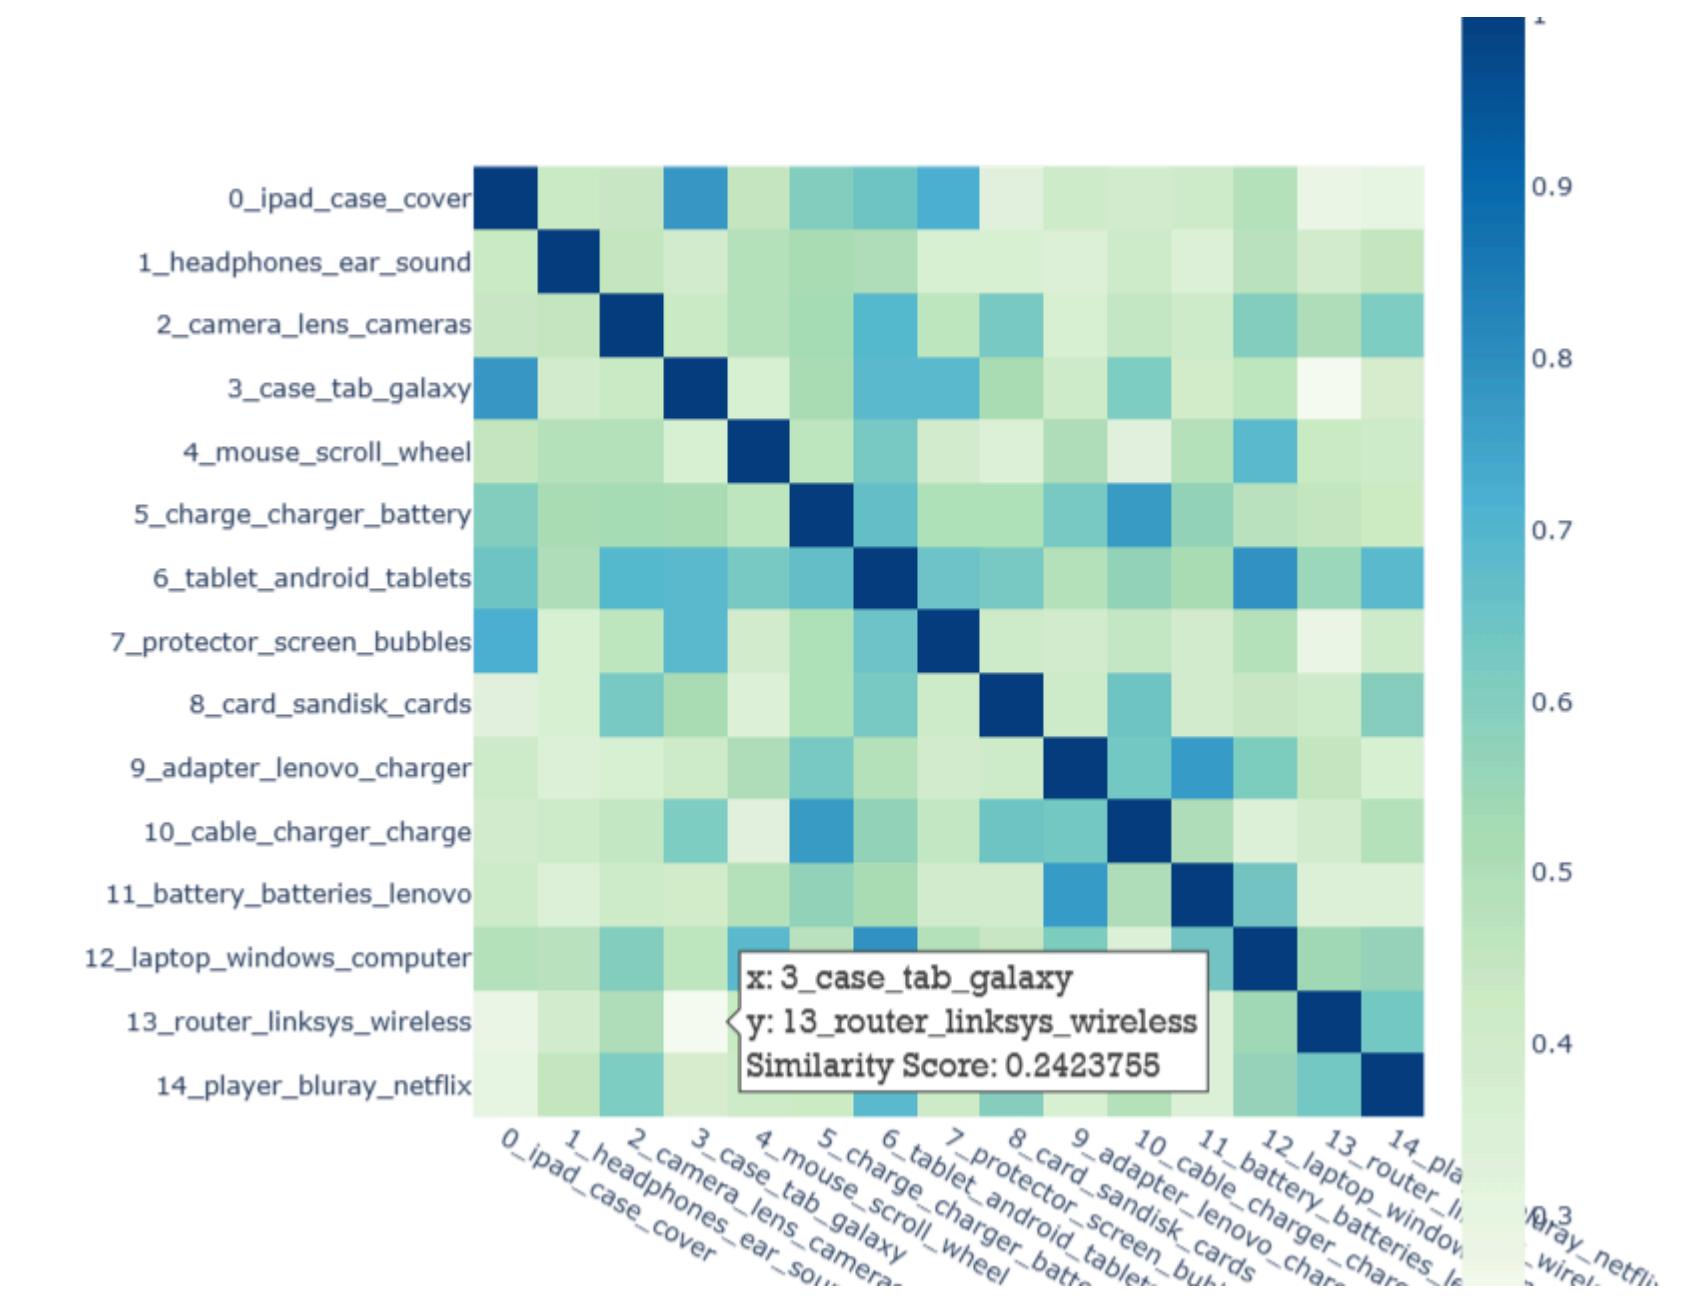
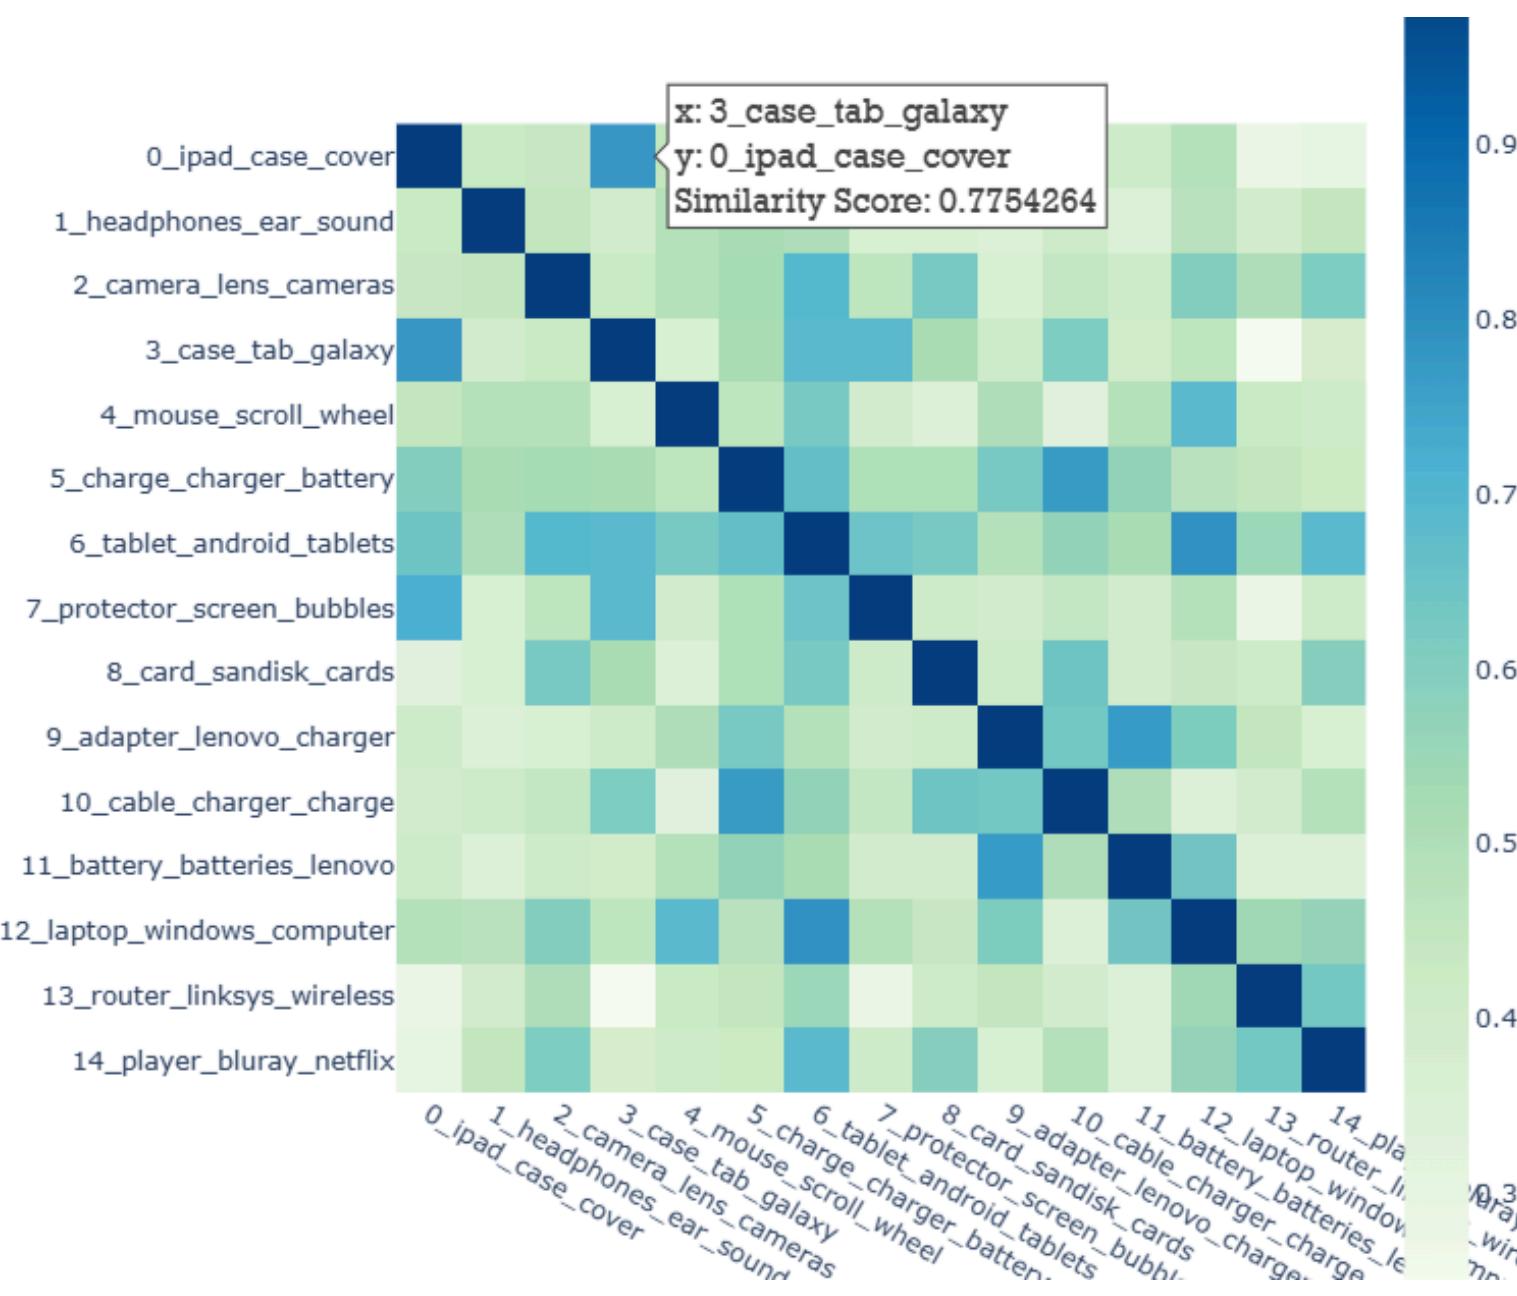


Results(similarity matrix)





Results(similarity matrix)

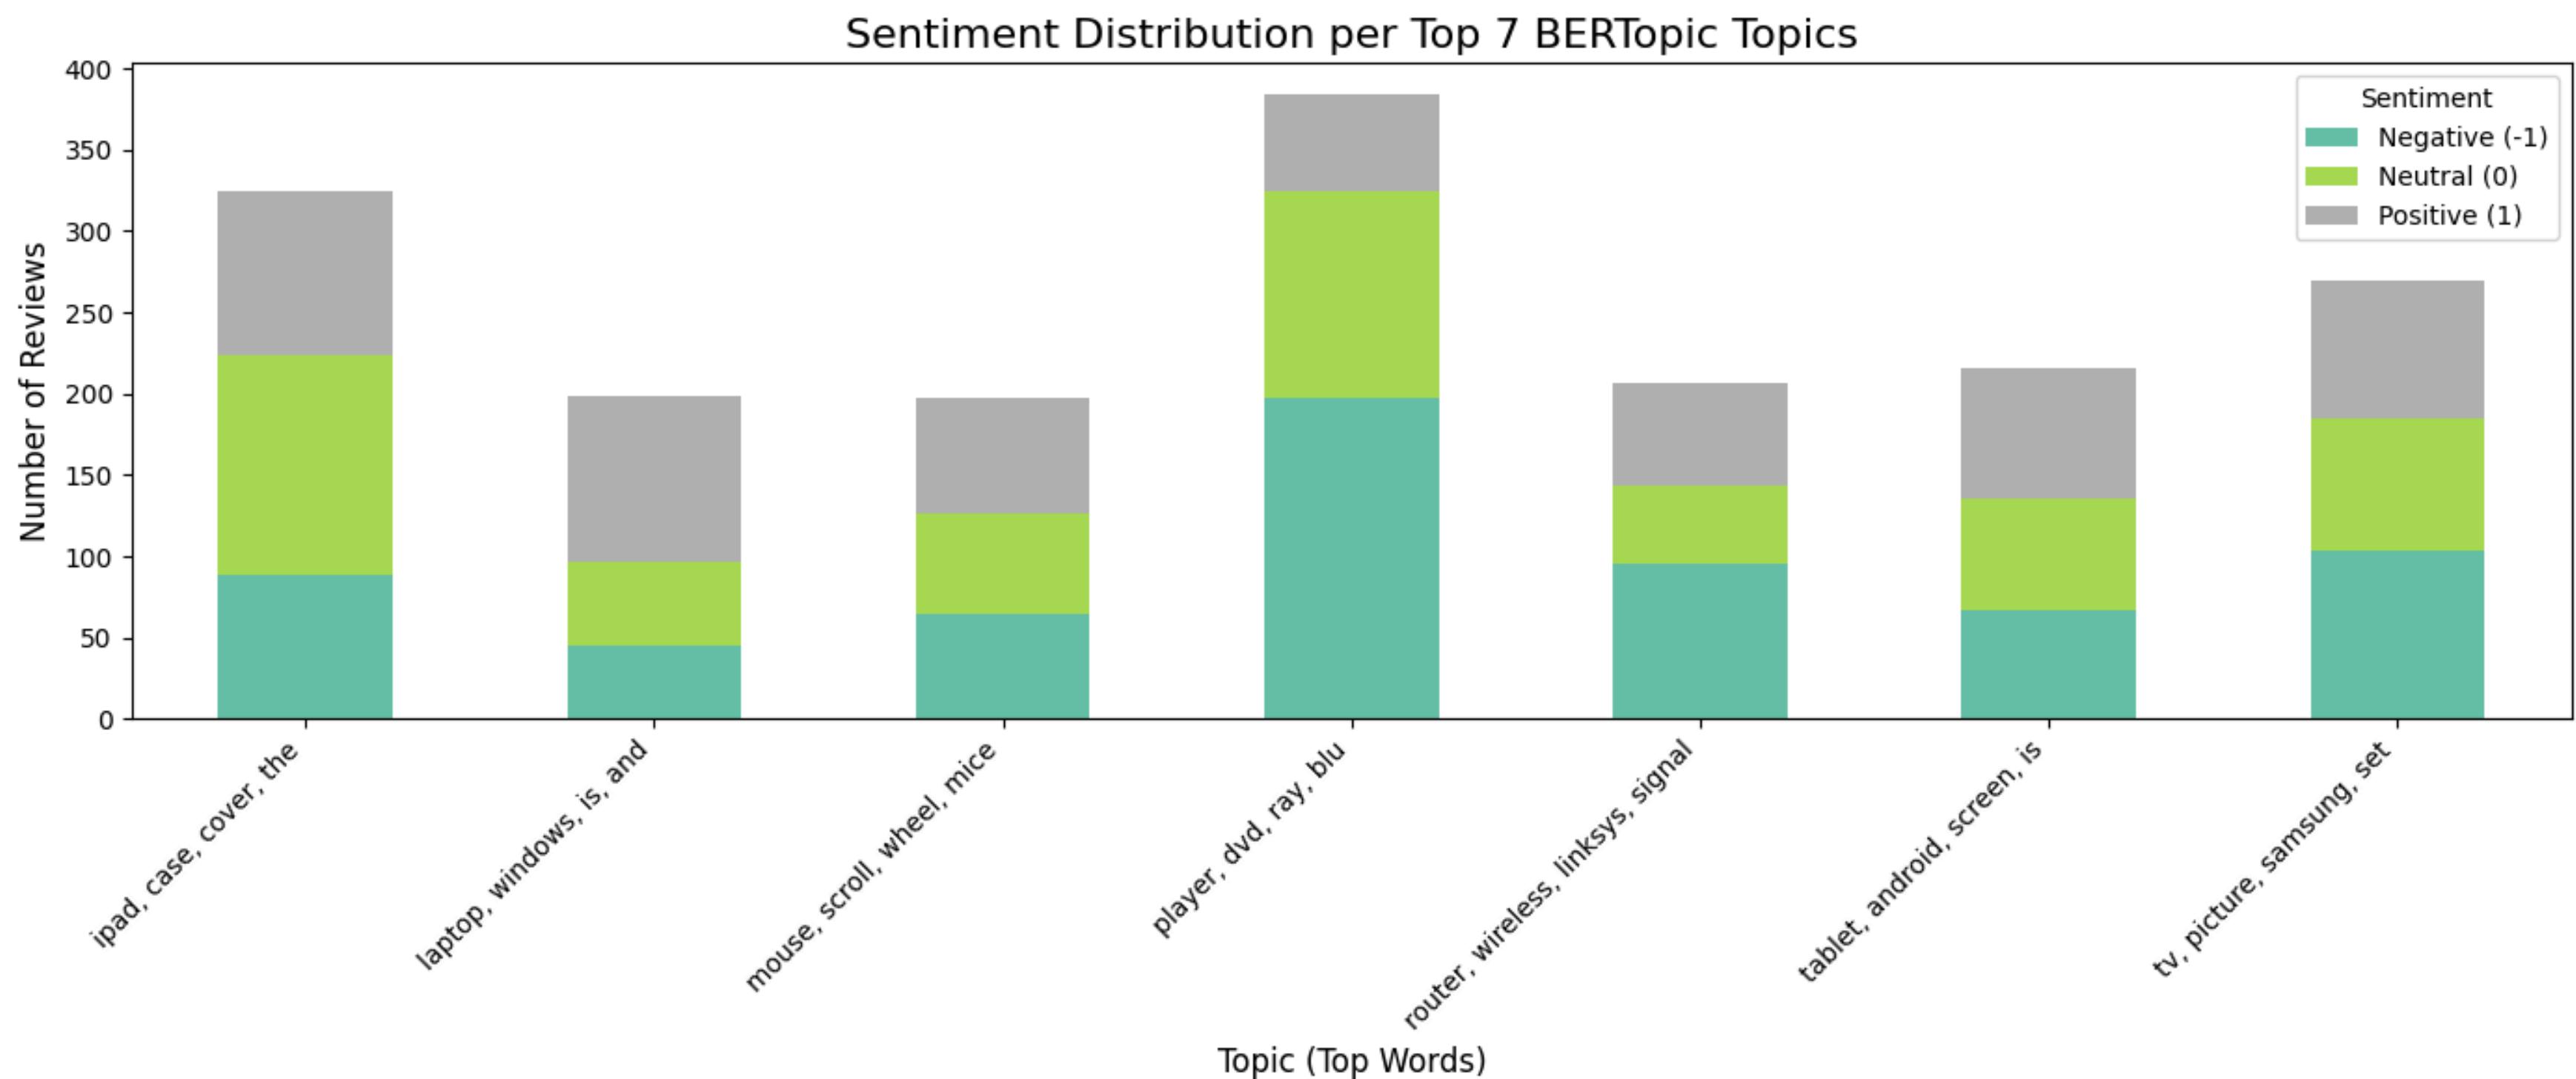


Results



- BERTopic utilizes transformer-based embeddings and clustering techniques to identify topics with better semantic coherence.
- Compared to LSA and NMF, BERTopic effectively differentiates between topics while maintaining interpretability.
- Topic sizes are not too huge, less overlapping

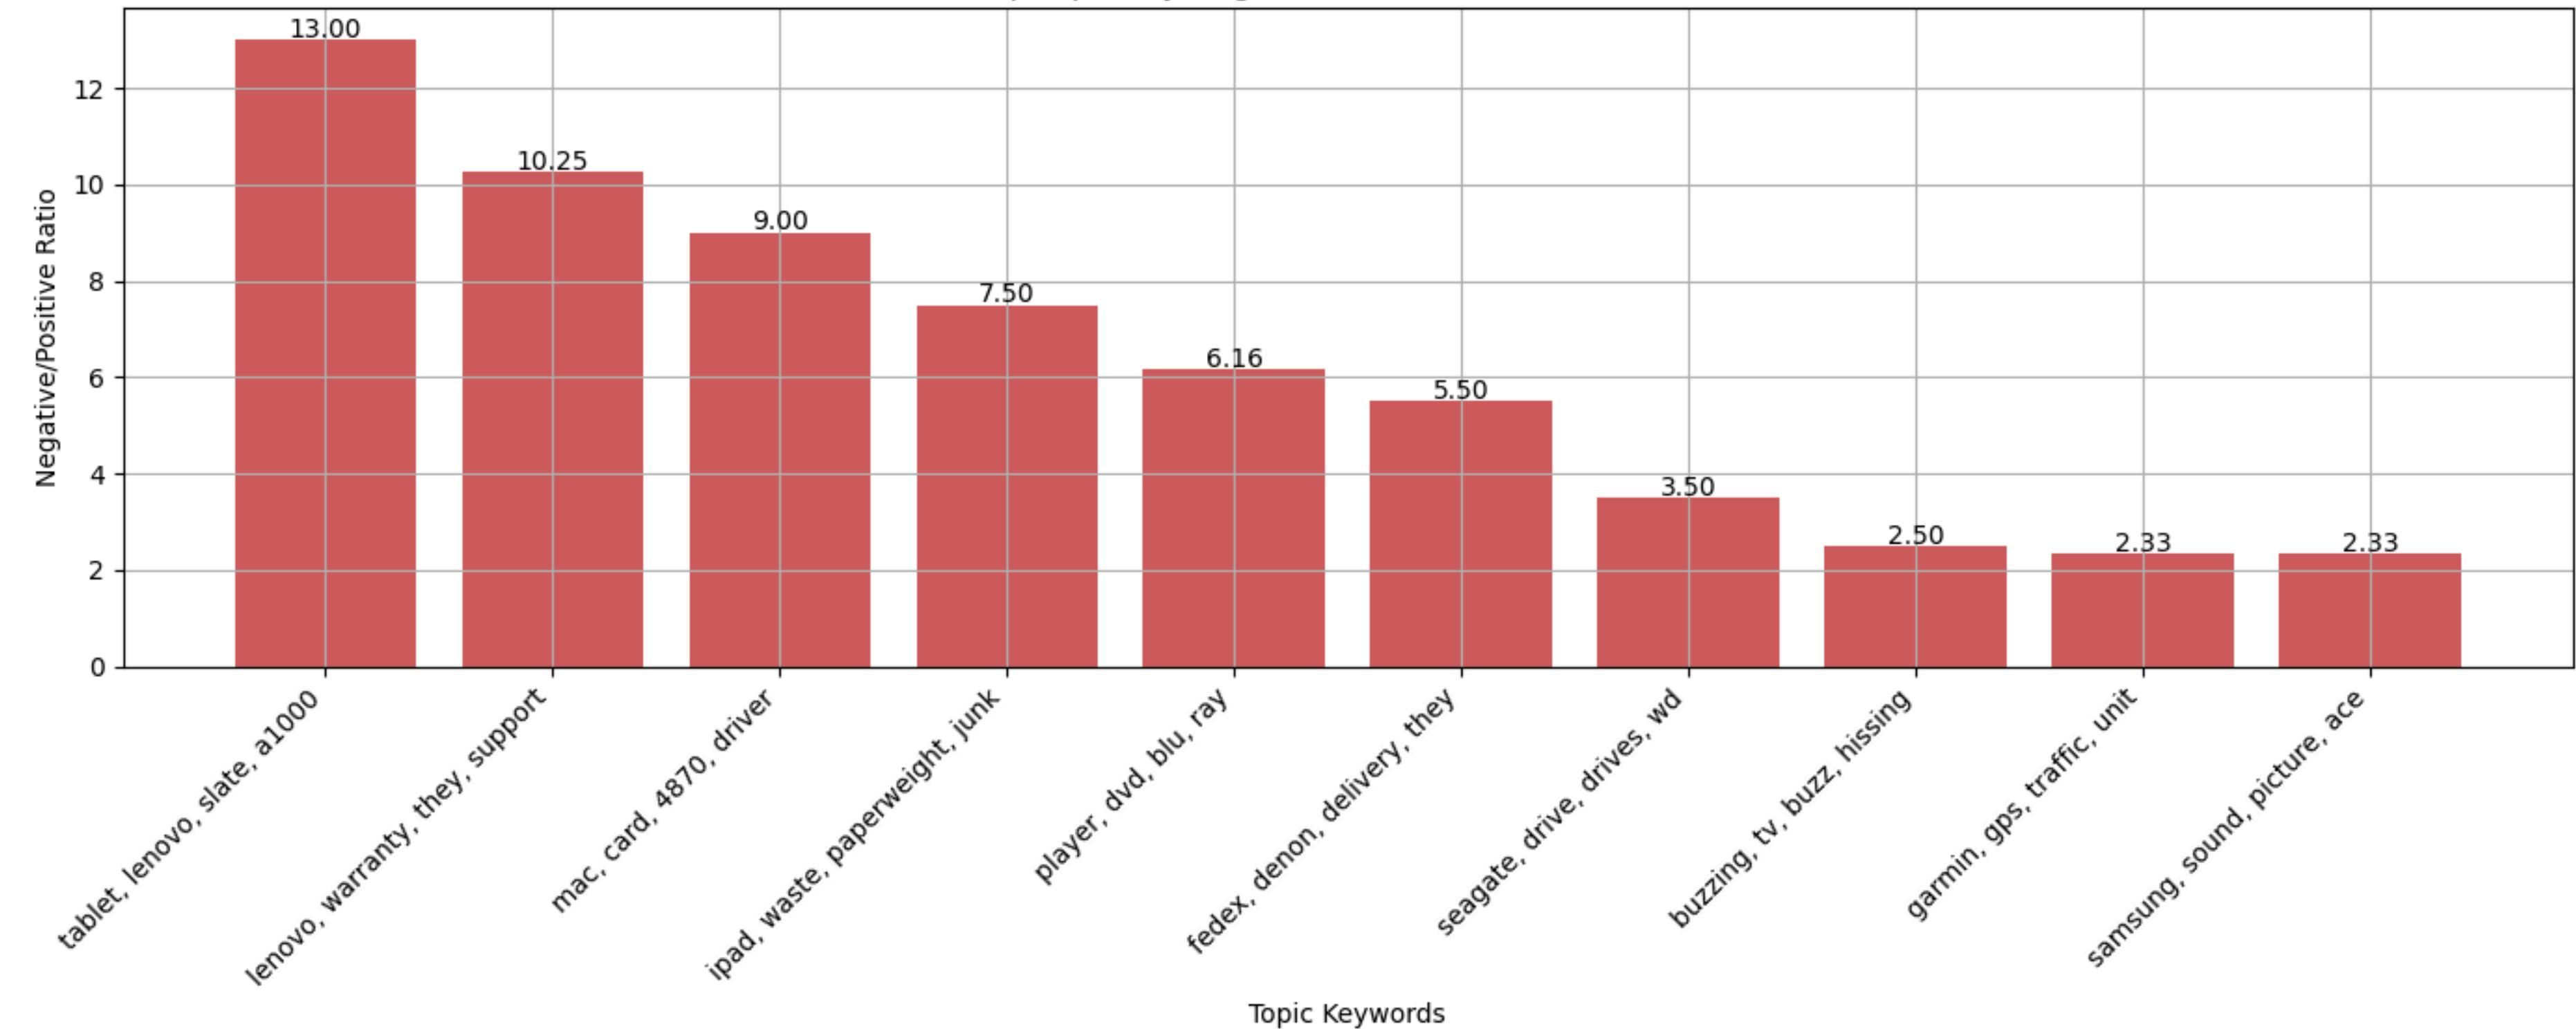
Sentiment Distribution Among Topics



Sentiment Distribution Among Topics



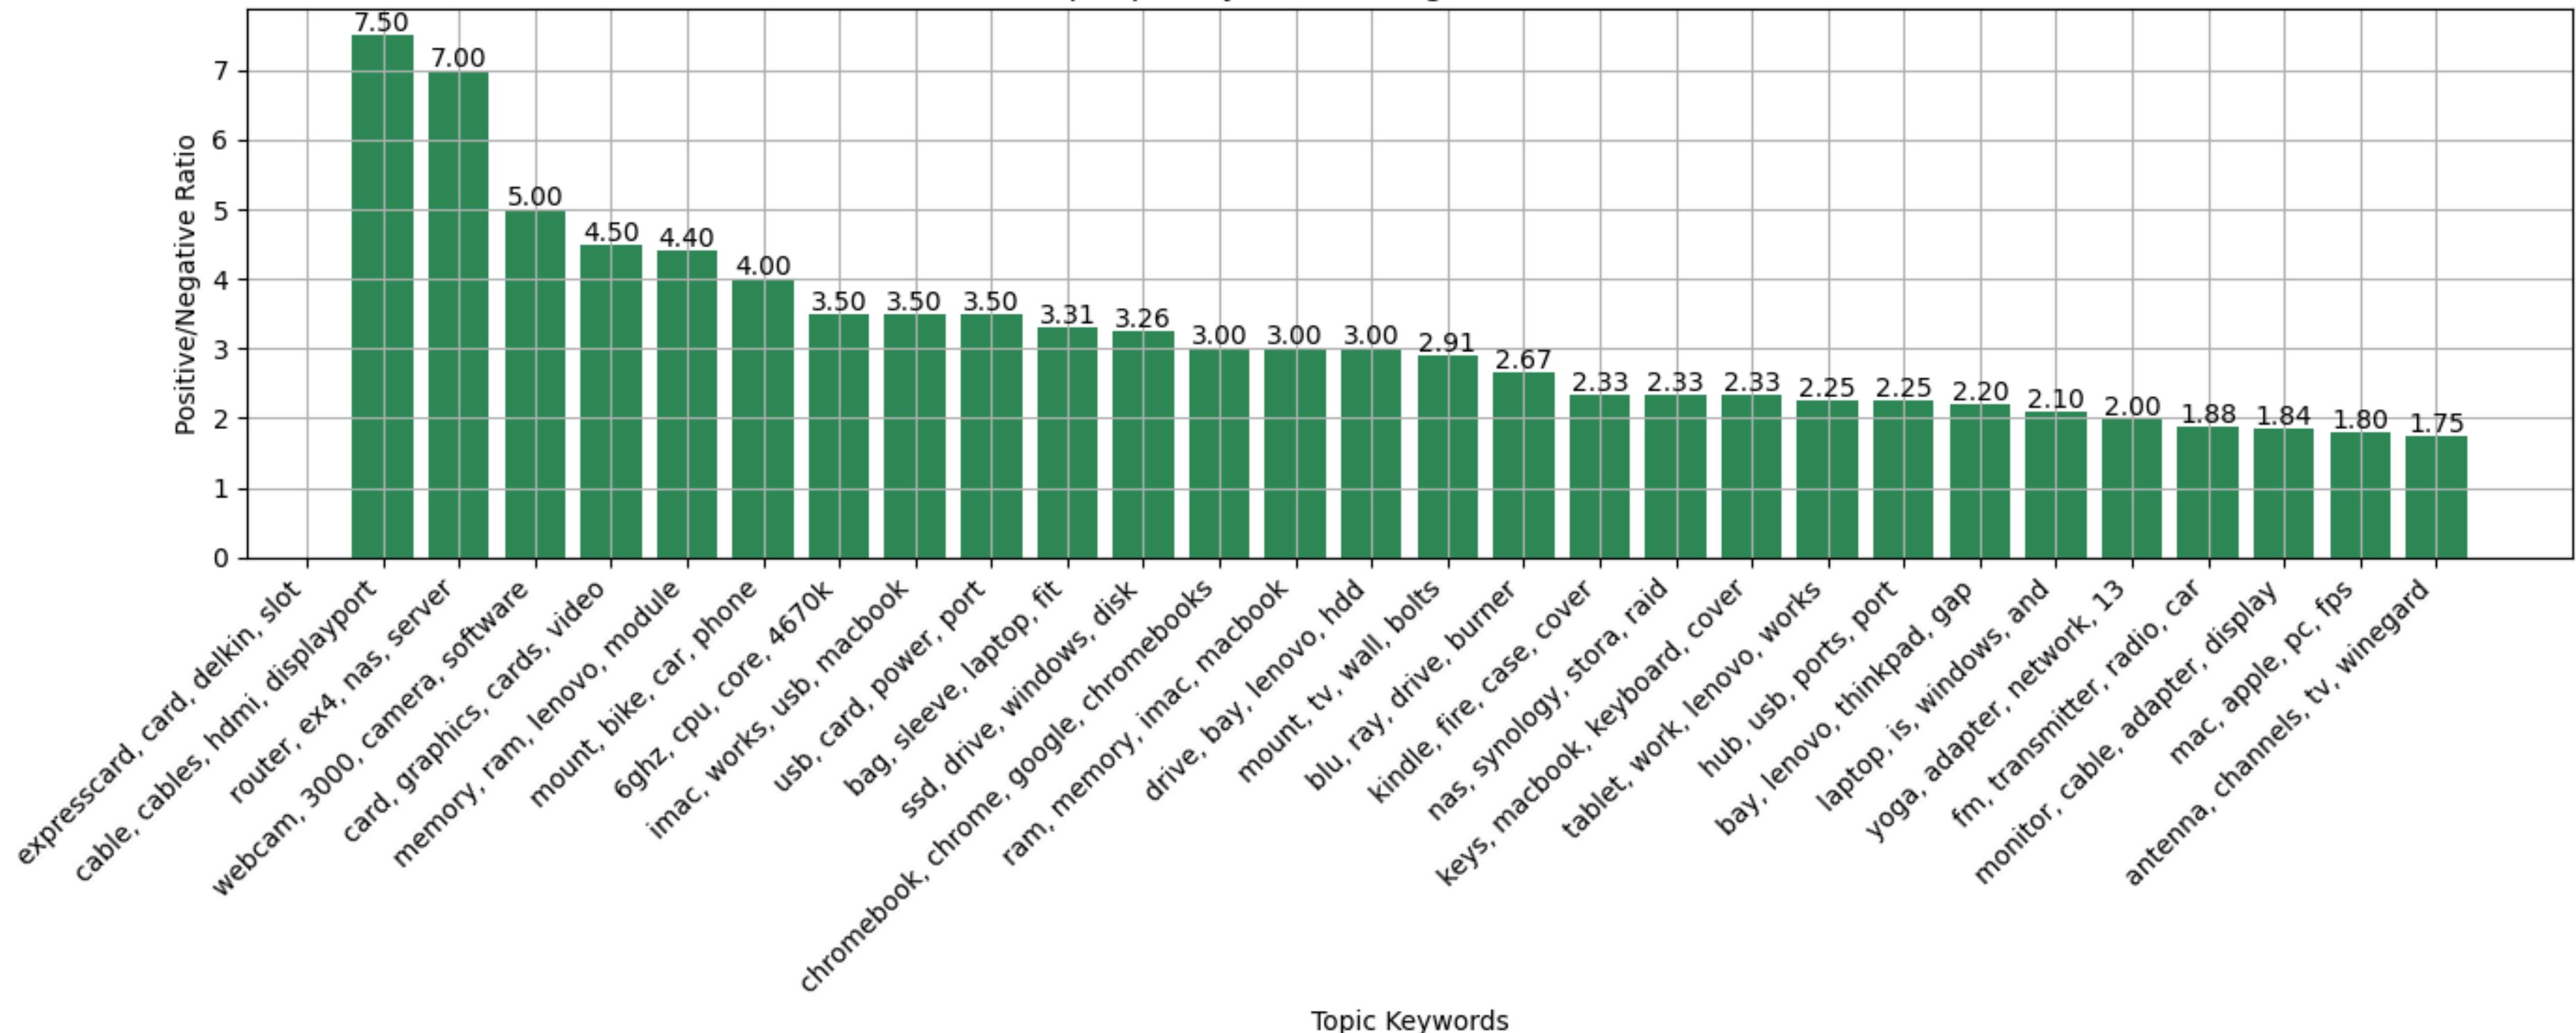
Top Topics by Negative/Positive Sentiment Ratio



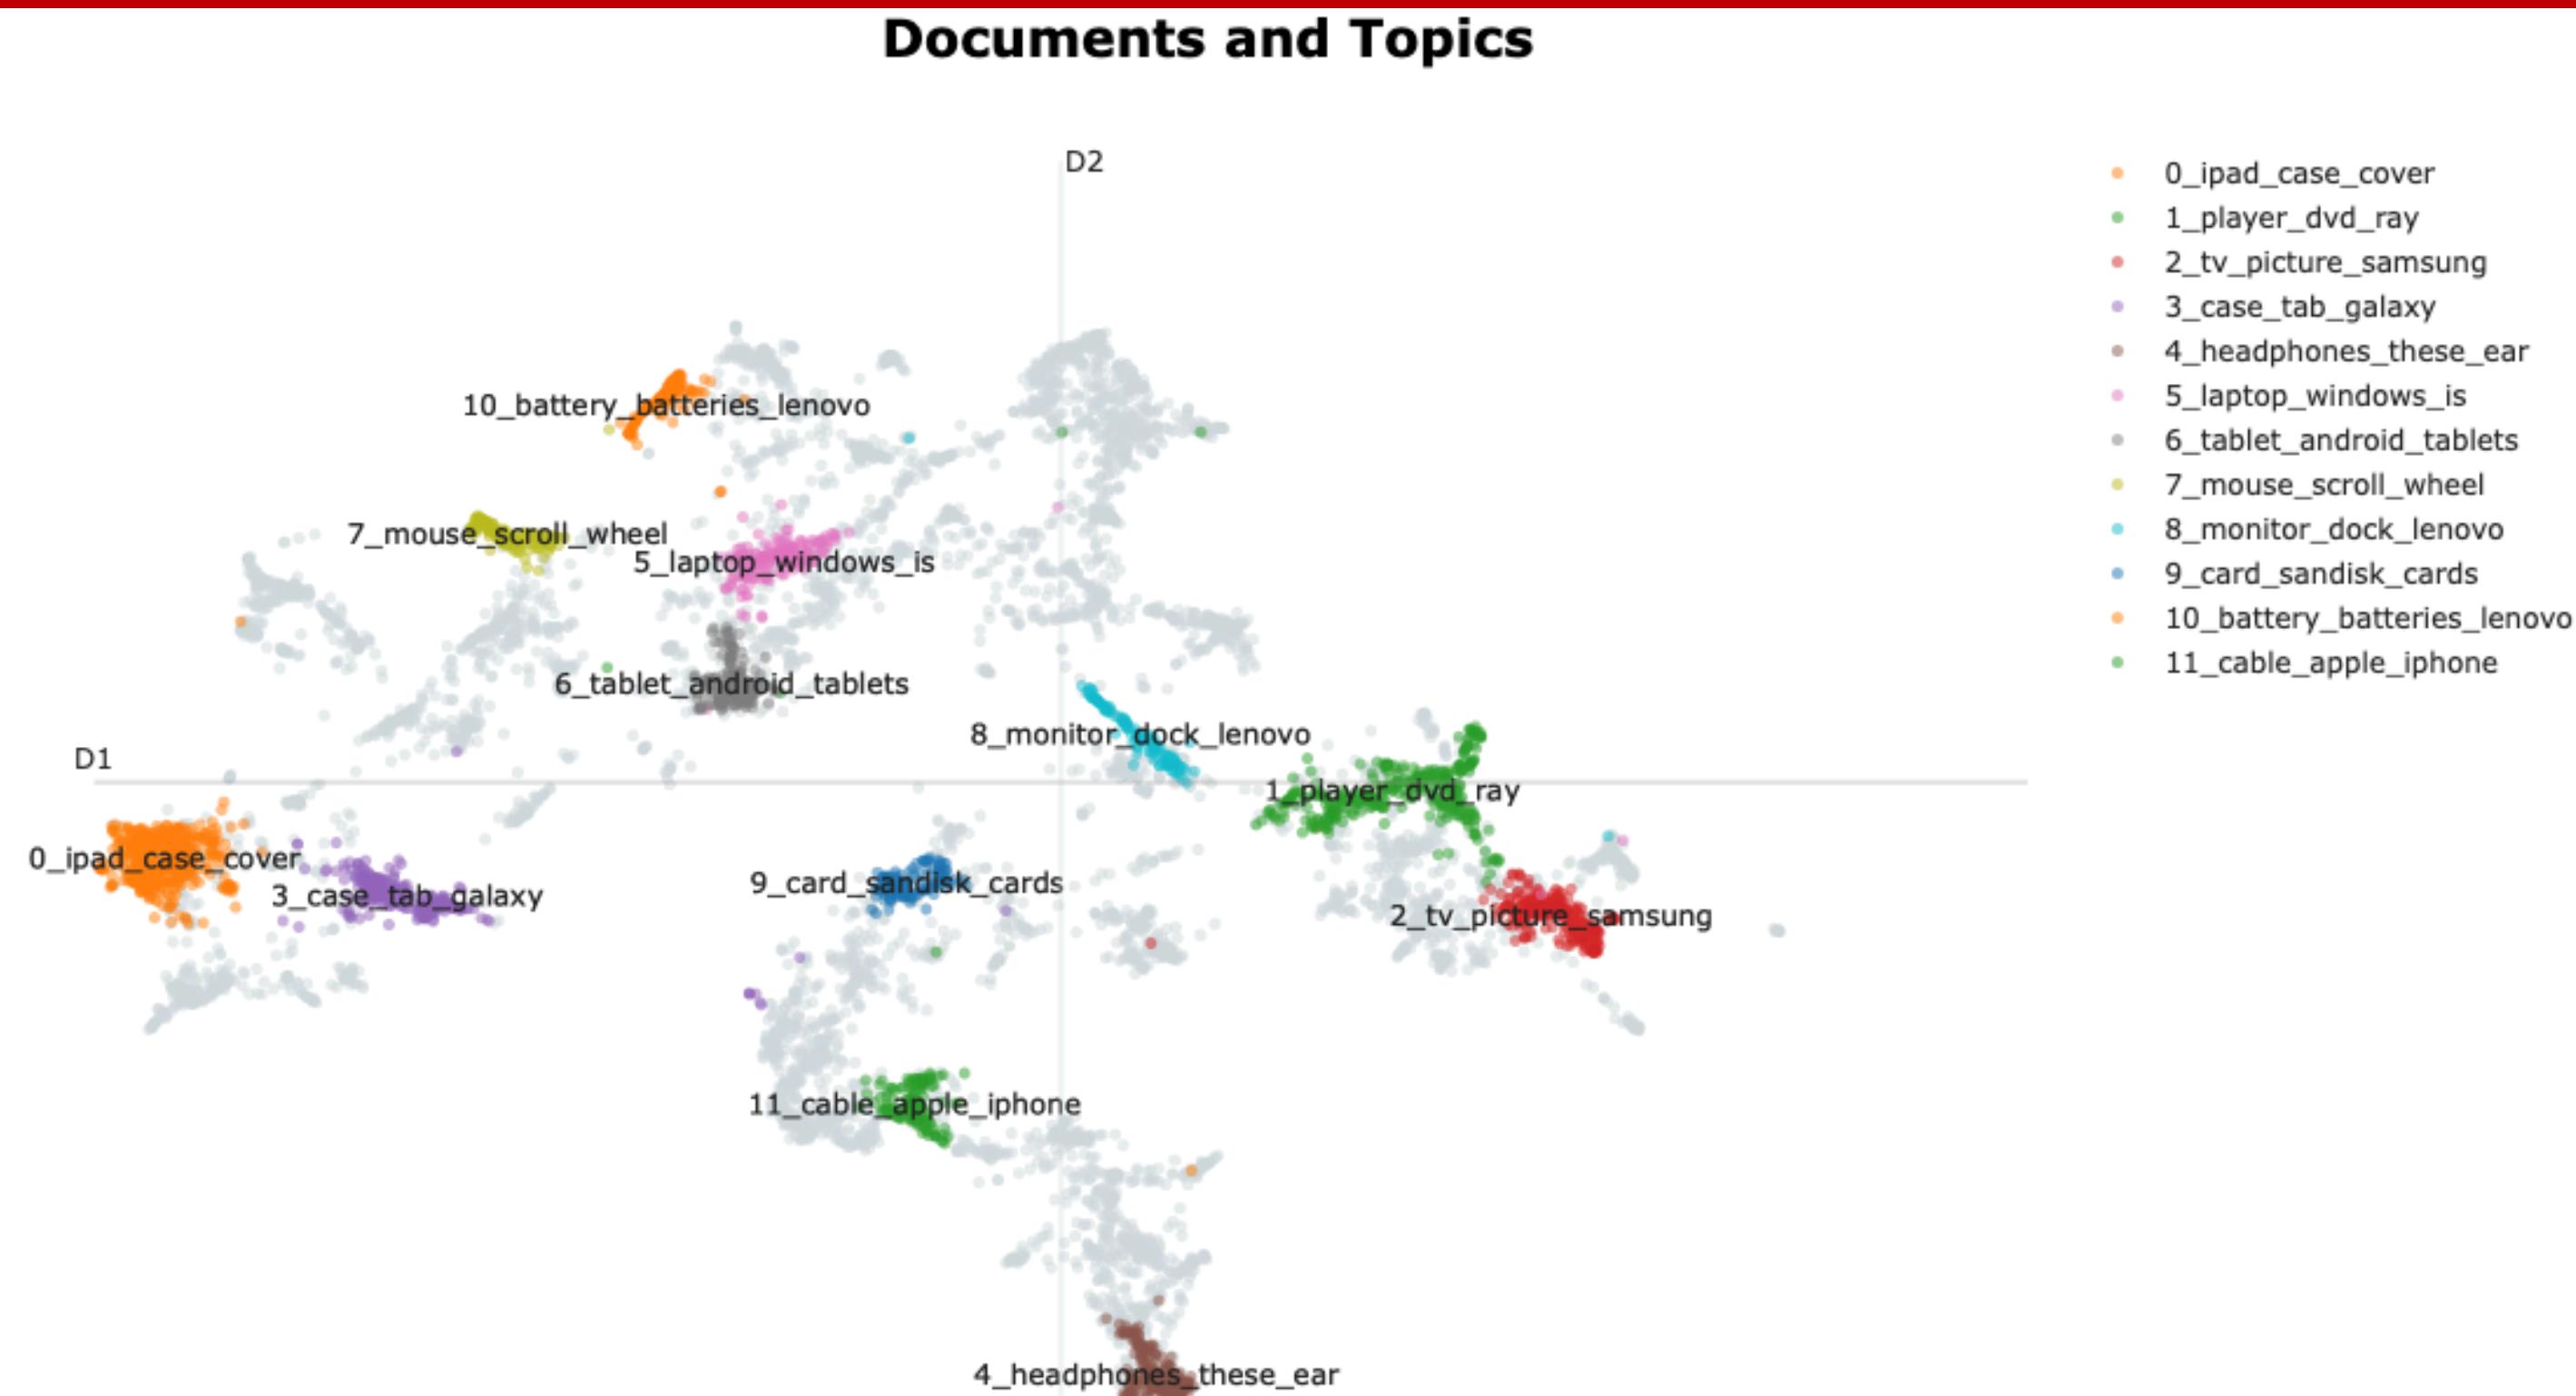
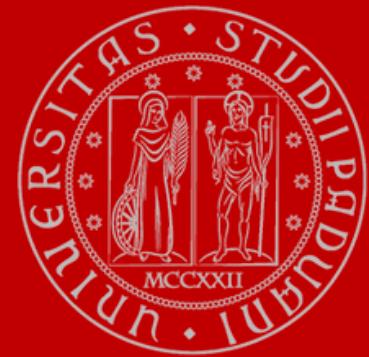
Semantic Distribution Among Topics



Top Topics by Positive/Negative Sentiment Ratio



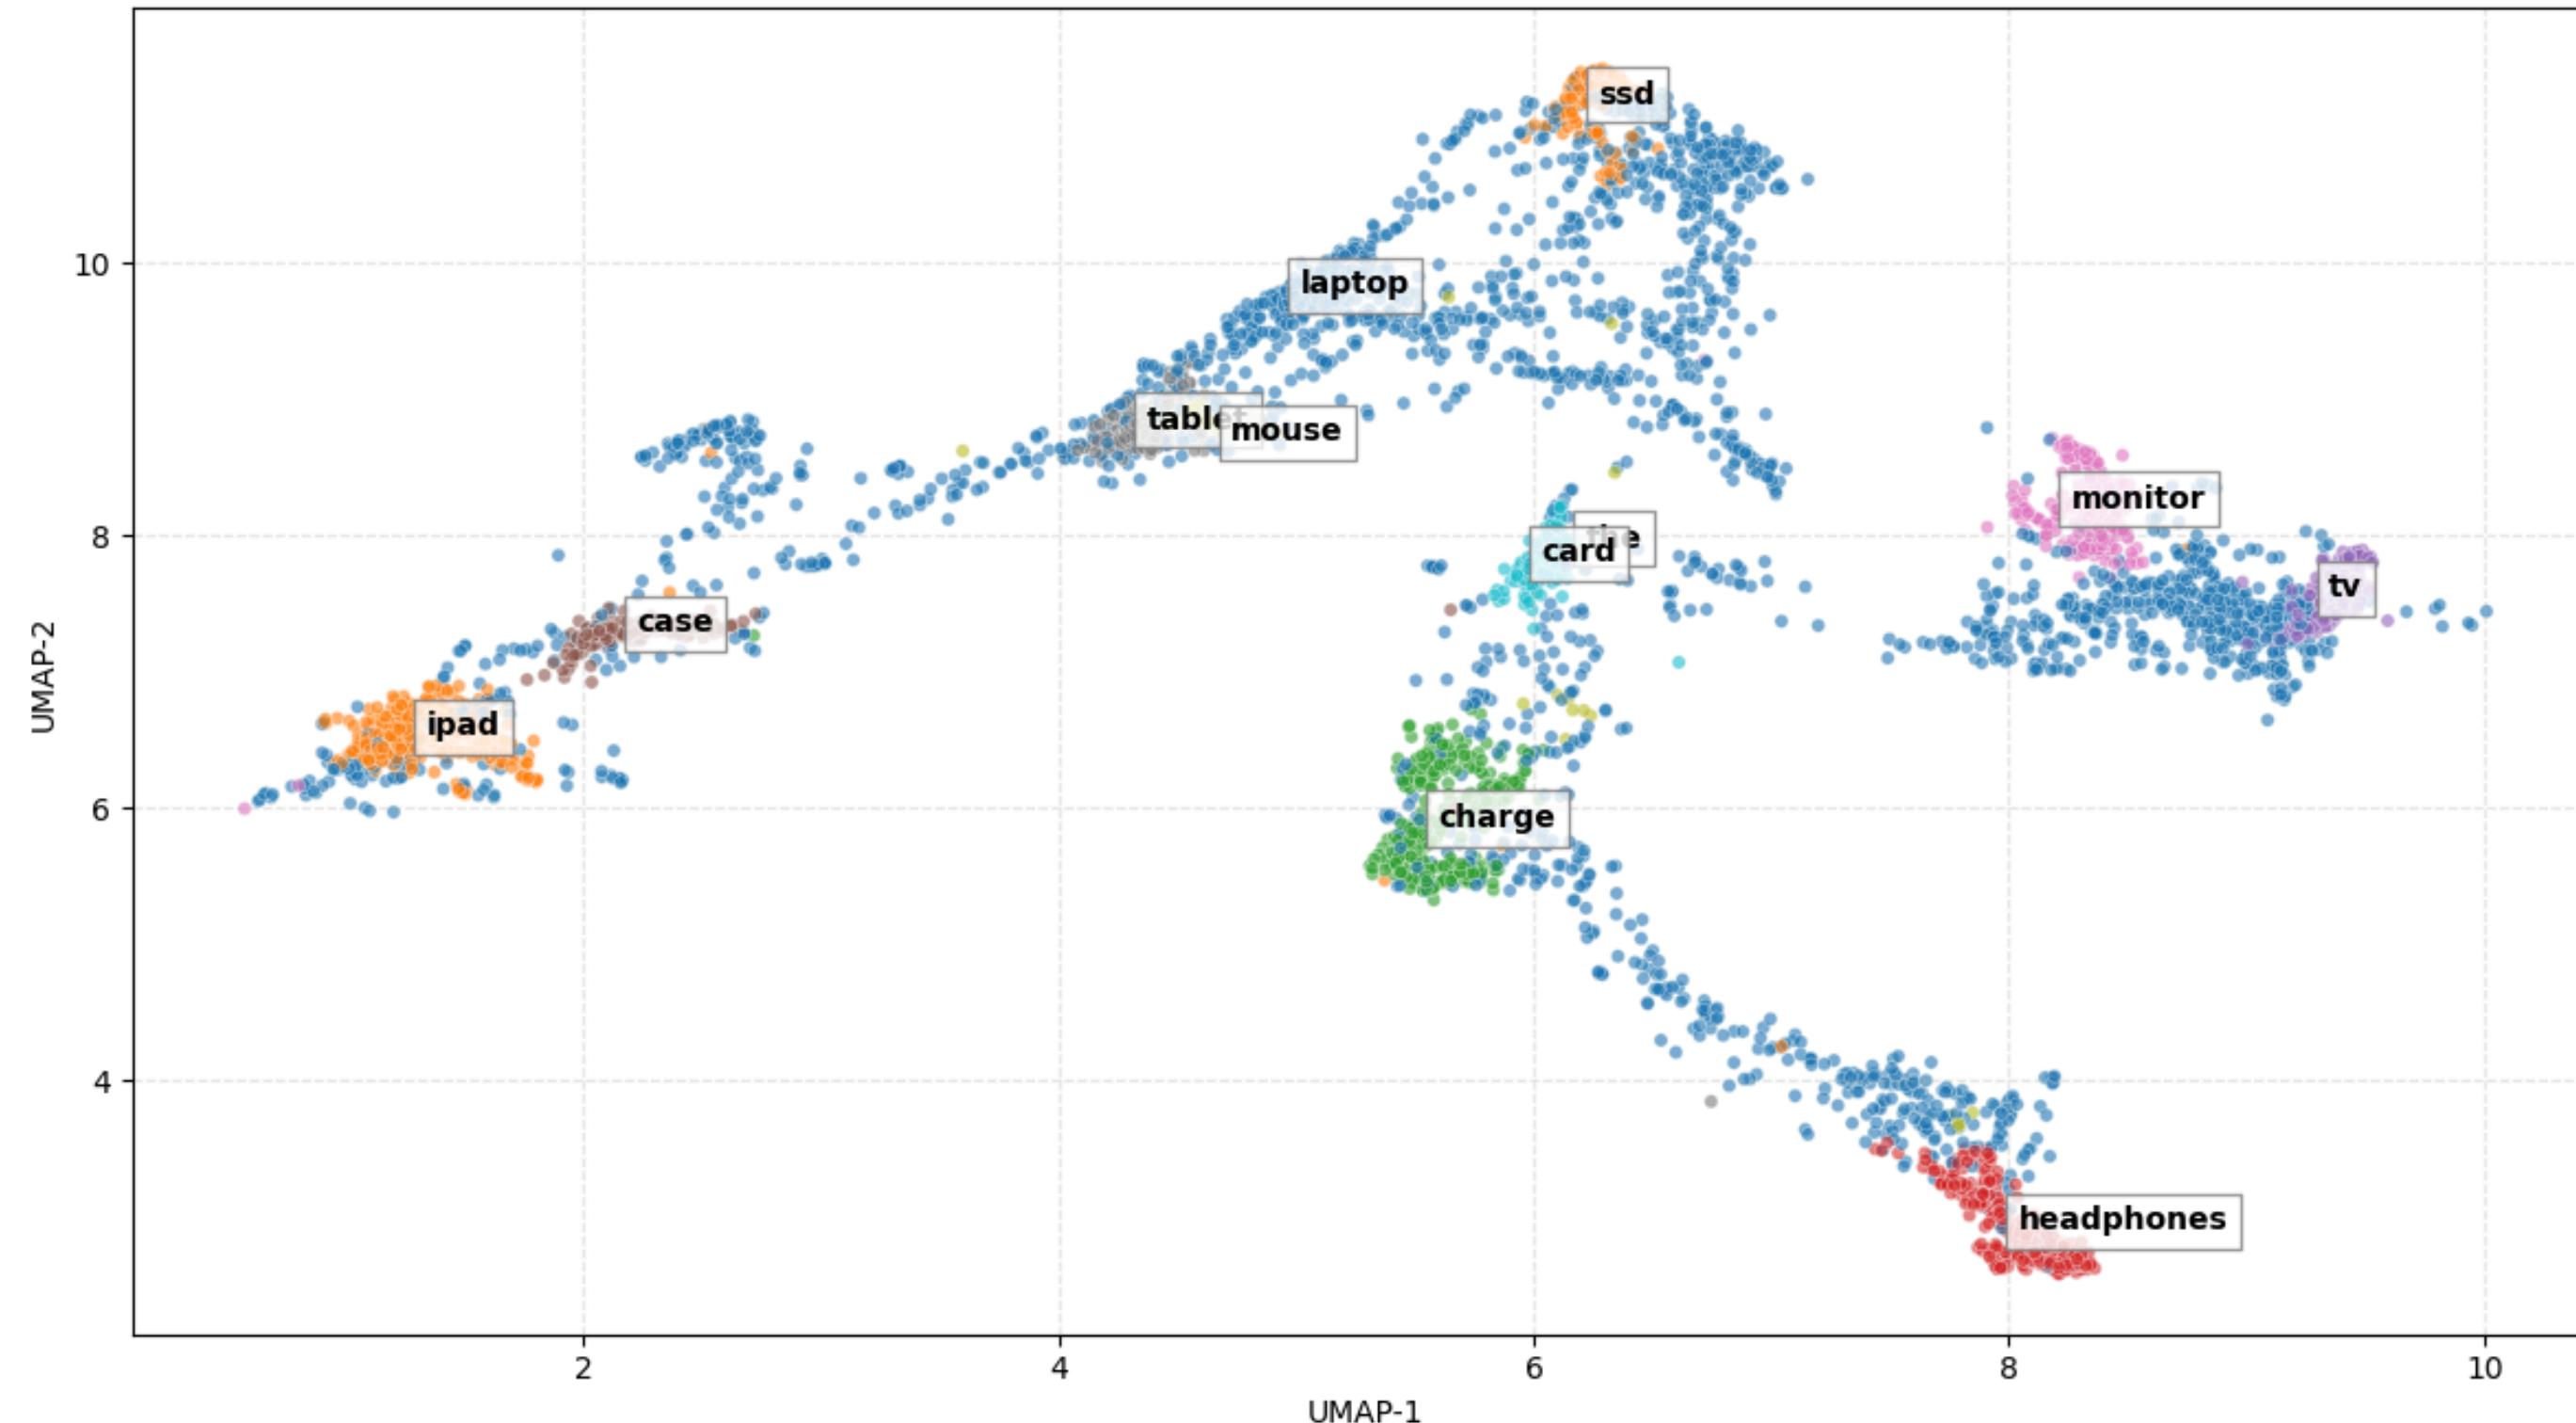
Sentiment Distribution Among Topics (continued)



results



Top 12 BERTopic Clusters with Representative Keywords



Conslusions



- Among the tested topic modeling methods, BERTopic demonstrated the best clustering performance, producing more semantically coherent topics. Traditional techniques like LSA and NMF struggled with excessive topic overlap, rendering their clusters less useful. The improved performance of BERTopic can be attributed to its embeddingbased approach, which captures deeper semantic relationships in textual data.

contributions



- Mohammad : data crawling - data cleaning and preparation - labeling - a transformers fine-tuning - fasttext - LSTM,CNN and classic models - writing reports
- Bahar: data labeling - data cleaning and preparation - a transformers fine-tuning - topic modelling and Semantic Distribution Among Topics - writing reports-fasttext

