

# Semantic Analysis

Mohammadkazem Rajabi

Mohammadkazem.rajabi@unipd.studenti.it

Baharehsadat Khatami

Baharehsadat.khatami@unipd.studenti.it

## 1. Introduction

Understanding customer sentiment in product reviews is a critical task for companies aiming to assess brand perception and improve user experience. In this project, we explore sentiment classification on user reviews of popular consumer technology products—primarily from Apple and Samsung—using a range of traditional and state-of-the-art NLP methods. The goal is to classify each review as expressing a positive, neutral, or negative sentiment.

While Transformer-based models have become the standard for many NLP tasks, we aim to systematically evaluate how much performance gain they provide compared to simpler and more efficient models such as FastText and recurrent neural networks. We also investigate whether traditional models, when enhanced with attention mechanisms and advanced loss functions like focal loss, can remain competitive.

Our experimental setup includes FastText[10] as a baseline, multiple deep learning models including LSTM[9], GRU[4], and BiLSTM[6] with attention mechanisms, and Transformer-based architectures such as RoBERTa[12], XLNet [13], and DeBERTa-v3[8]. In addition, we apply stratified K-fold cross-validation for robust evaluation, and focal loss to address class imbalance during training.

Beyond supervised classification, we incorporate topic modeling to explore latent semantic themes within the review corpus. We apply four complementary techniques: Latent Dirichlet Allocation (LDA)[3], Latent Semantic Analysis (LSA)[5], Non-negative Matrix Factorization (NMF)[11], and BERTopic[7]. These models help uncover brand-specific thematic structures and offer interpretability that complements the results of our sentiment classification pipeline.

The results show that DeBERTa-v3 consistently outperforms other models in terms of both accuracy and F1-score, while attention-based[1] BiLSTM models demonstrate competitive performance with significantly lower computational cost. Topic modeling further enhances our understanding of the underlying semantic space, providing a holistic view of sentiment and themes in real-world, domain-specific datasets.

## 2. Dataset

The dataset used in this project is derived from the publicly available Amazon Electronics 5-core dataset provided by Stanford SNAP. It contains millions of user reviews across a wide range of electronic products. We programmatically filtered this dataset to extract reviews specifically related to Apple, Samsung, and Lenovo products.

Brand identification was performed using a set of manually curated keyword lists. Reviews were matched to a brand (Apple, Samsung, or Lenovo) if they contained relevant product keywords such as “iPhone,” “Galaxy,” or “ThinkPad,” while reviews containing irrelevant topics (e.g., “pizza,” “restaurant”) were excluded through a blacklist filtering process.

Each review includes a star rating from 1 to 5. We converted these ratings into sentiment labels using the following mapping: 1–2 stars as **negative** sentiment (-1), 3 stars as **neutral** sentiment (0), and 4–5 stars as **positive** sentiment (1). This produced a labeled sentiment dataset grounded in real user ratings.

To improve neutral class representation—often under-represented in natural distributions—we supplemented the data with a set of synthetic neutral reviews generated using large language models (LLMs). These synthetic samples were created through prompt engineering to elicit objective, descriptive statements that reflect neutral sentiment. The final dataset contains approximately 17,000 labeled reviews, covering a balanced mix of brands and sentiment classes.

## 3. Methodology

### 3.1. Preprocessing

To ensure consistency and reduce noise, we applied a sequence of standard NLP preprocessing techniques to all input texts. These included lowercasing, punctuation removal, and tokenization. Stopwords were removed using the NLTK corpus to reduce sparsity and emphasize semantically meaningful tokens. We deliberately avoided stemming and lemmatization to retain lexical nuance, which is beneficial for pre-trained Transformer-based models.

## 3.2. Model Architectures

### 3.2.1 FastText Baseline

We adopted FastText as a lightweight baseline to establish reference performance. FastText represents input sequences as averages of word embeddings and incorporates subword-level information, enabling efficient handling of rare and misspelled words. This model is particularly suitable for large-scale or real-time applications due to its low computational footprint.

### 3.2.2 Deep Neural Models

We implemented several recurrent neural network (RNN) architectures, including unidirectional LSTM, GRU, and bidirectional LSTM (BiLSTM). These architectures model temporal dependencies in sequential text data and are known to capture syntactic and semantic patterns effectively. To improve contextual representation, we extended the BiLSTM model with a self-learned attention mechanism. The attention mechanism computes a weighted aggregation over hidden states, enabling the model to focus on sentiment-relevant tokens dynamically.

All RNN-based models utilized randomly initialized embedding layers, followed by recurrent layers with hidden dimension sizes between 64 and 128. A dropout layer was applied after the recurrent units to mitigate overfitting, followed by a fully connected output layer with softmax activation for multi-class classification.

### 3.2.3 Transformer-based Models

We fine-tuned three Transformer-based language models—RoBERTa, XLNet, and DeBERTa-v3—using the HuggingFace Transformers library. These models are pretrained on massive corpora using masked or permutation-based language modeling objectives and have demonstrated state-of-the-art performance on numerous NLP tasks.

RoBERTa improves over BERT by eliminating the next sentence prediction task and increasing training efficiency through dynamic masking and longer sequences. XLNet introduces a permutation language modeling objective that captures bidirectional context while maintaining autoregressive capabilities. DeBERTa-v3 employs disentangled attention mechanisms and relative position encodings, offering superior contextual understanding, especially in low-resource settings.

For each model, we appended a linear classification head to the final encoder layer and fine-tuned the entire network and Early stopping was used based on validation loss.

## 3.3. Training Strategy and Evaluation

To address class imbalance, particularly the underrepresented neutral class, we employed the focal loss function.

Focal loss dynamically scales the loss contribution of each instance based on classification difficulty, thereby focusing training on harder, misclassified examples.

All models were trained and evaluated using stratified 5-fold cross-validation to ensure balanced representation of sentiment classes across splits. We used macro-averaged metrics to assess model performance, including accuracy, precision, recall, and F1-score. Confusion matrices were also generated for qualitative analysis of class-wise predictions and misclassification patterns.

## 3.4. Transformer Model Performance

To evaluate the effectiveness of pre-trained Transformer architectures in sentiment classification, we fine-tuned six language models: DeBERTa-v3, DeBERTa-v3-base, DeBERTa-v3-large, RoBERTa, XLNet, and CardiffNLP[2]. These models have shown strong performance across various NLP tasks, and we assess their suitability for multi-class sentiment classification on our dataset of product reviews.

All models were fine-tuned using a linear classification head on top of the final hidden state representation, with training carried out using early stopping and the AdamW optimizer. We evaluated the models using macro-averaged metrics on a stratified held-out test set. The results are reported in Table 1.

Table 1. Test set performance of Transformer-based models.

Model	Accuracy	Precision	Recall	F1-score
CardiffNLP	<b>0.7598</b>	0.7524	<b>0.7598</b>	<b>0.7546</b>
DeBERTa-v3-large	0.7526	<b>0.7539</b>	0.7526	0.7532
DeBERTa-v3-base	0.7228	0.7572	0.7228	0.7332
DeBERTa-v3	0.7291	0.7266	0.7291	0.7277
RoBERTa	0.7260	0.7204	0.7260	0.7223
XLNet	0.7088	0.7371	0.7088	0.7180

The results show that CardiffNLP achieves the highest overall performance across all metrics, closely followed by DeBERTa-v3-large. Despite being smaller in size, DeBERTa-v3-base exhibits strong precision, suggesting it is a viable option when computational efficiency is a priority. RoBERTa and XLNet deliver reasonable performance but lag behind the DeBERTa family. Interestingly, XLNet achieves high precision but struggles with recall, possibly due to overfitting or instability during fine-tuning. These results reinforce the effectiveness of modern Transformers for sentiment classification, particularly those using disentangled attention mechanisms and large-scale pretraining.

## 3.5. Deep Learning Model Performance

We implemented and compared a range of deep learning models to perform sentiment classification on our multi-class dataset. These include baseline BiLSTM architectures as well as enhanced models using convolutional layers, at-

tention mechanisms, pretrained word embeddings (GloVe), and focal loss to address class imbalance.

All models were trained on the same dataset using either stratified train-test splits or 5-fold cross-validation (CV), and evaluation was based on test accuracy and macro-averaged F1-score.

Table 2. Test set performance of deep learning models.

Model	Accuracy	F1-score
<b>BiLSTM + Attention + GloVe</b>	<b>0.7000</b>	0.6600
GRU + GloVe	<b>0.7084</b>	<b>0.7099</b>
BiLSTM + Attention	<b>0.7039</b>	0.6800
Conv1D + BiLSTM	0.6817	0.6820
Vanilla BiLSTM	0.6674	0.6696
GRU	0.6348	0.6408

**Model descriptions:** The **Vanilla BiLSTM** model is a baseline architecture consisting of two stacked Bidirectional LSTM layers with dropout and dense output. The **BiLSTM + Attention** model enhances this by adding a self-attention mechanism between the LSTM layers to help the model focus on informative tokens. The **GRU + GloVe** model uses pretrained GloVe word embeddings and replaces LSTM layers with GRU units, improving generalization. The **GRU** model uses GRU layers without GloVe and demonstrates weaker performance, emphasizing the benefit of pretrained embeddings.

The **Conv1D-BiLSTM** model introduces a 1D convolutional layer before the BiLSTM to extract local n-gram patterns, followed by global max pooling and classification layers. Lastly, the **BiLSTM + Attention + GloVe + Focal Loss** model combines the strengths of pretrained embeddings, attention mechanisms, and focal loss to effectively handle class imbalance and delivers the most balanced overall performance across all classes.

### 3.6. FastText Sentiment Classification

We also explored FastText, a lightweight and efficient text classification algorithm developed by Facebook AI [?]. FastText performs classification using bag-of-words and subword information, offering a highly scalable alternative to more resource-intensive neural models.

**Preprocessing and Format:** We used the same dataset of 17,000 reviews labeled across three sentiment classes (negative, neutral, positive). Reviews were converted to FastText format by prefixing each entry with `__label__NEG`, `__label__NEU`, or `__label__POS`.

**Training and Evaluation:** The dataset was split using a stratified 80/20 train-test split. The model was trained using FastText’s supervised mode with default parameters. Evaluation was performed on the test set using Scikit-learn, and the trained model was approximately 895MB.

#### Overall Evaluation:

Table 3. Class-wise performance of FastText model

Class	Precision	Recall	F1-score
Negative (NEG)	0.6750	0.5807	0.6243
Neutral (NEU)	0.4803	0.4699	0.4751
Positive (POS)	0.7444	0.8177	0.7794

- **Macro-Averaged F1-score:** 0.6263

- **Overall Accuracy:** 66.15%

**Discussion:** FastText achieved solid performance considering its speed and simplicity. The model was most confident in predicting positive sentiment, while neutral classification remained the most challenging—consistent with our findings in deep learning and Transformer models. Its efficiency and small training footprint make it an attractive option for large-scale or resource-constrained deployment.

### 3.7. Shallow Neural Networks and Classical ML Baselines

To complement our evaluation of advanced deep learning models, we also examined the performance of shallow neural networks and classical machine learning baselines. These included simplified GRU-based architectures and traditional classifiers using TF-IDF features. While the shallow models demonstrated modest gains from bidirectionality, classical approaches generally underperformed. Overall, this comparison highlights the limitations of shallow and feature-based methods in capturing the nuanced structure of sentiment-laden text.

As shown in Table 4, shallow neural models and classical methods underperform compared to deeper architectures.

Table 4. Performance of shallow neural networks and classical ML baselines.

Model	Accuracy	F1
GRU (1 layer)	0.62	0.58
BiGRU (shallow)	0.66	64
<i>Logistic Regression (TF-IDF)</i>	0.59	0.55
<i>Random Forest and Logistic Regression</i>	0.64	0.58
<i>Decision Tree (TF-IDF)</i>	0.53	0.49

## 4. Topic Modeling Exploration

To enrich our understanding of the review dataset beyond sentiment polarity, we conducted an unsupervised topic modeling analysis. This allowed us to extract underlying themes, user concerns, and contextual patterns that frequently emerge in tech-related reviews. We employed four widely-used topic modeling approaches: Latent Dirichlet Allocation (LDA), Latent Semantic Analysis (LSA), Non-negative Matrix Factorization (NMF), and BERTopic.

#### 4.1. Preprocessing and Input Representation

All reviews were preprocessed by converting to lower-case, removing punctuation, and filtering out English stop-words. Lemmatization was not applied.

Depending on the algorithm, two types of vectorization were used:

- **TF-IDF Matrix:** Used for LSA and NMF to emphasize rare but important terms.
- **Bag-of-Words (BoW):** Used for LDA and BERTopic to preserve raw term frequency.

#### 4.2. Latent Dirichlet Allocation (LDA)

LDA [3] is a generative probabilistic model that represents documents as mixtures of topics, where each topic is a distribution over words. In our analysis, we trained the LDA model with 10 topics and visualized the top keywords per topic. The results revealed coarse-grained themes such as: we report the top 10 most probable words per topic in Table 5.

Table 5. Top 5 keywords per topic extracted using LDA.

Topic	Top Keywords
Topic 1	cable, usb, adapter, sound, iphone
Topic 2	laptop, lenovo, windows, mouse, computer
Topic 3	drive, card, usb, gb, hard
Topic 4	support, router, wireless, device, wifi
Topic 5	keyboard, keys, key, typing, bluetooth
Topic 6	tv, samsung, sound, player, quality
Topic 7	one, would, battery, samsung, get
Topic 8	case, ipad, screen, like, cover

#### 4.3. Latent Semantic Analysis (LSA)

LSA [5] is a dimensionality reduction technique based on Singular Value Decomposition (SVD) of the TF-IDF matrix. LSA captures linear substructures and latent relationships in term usage. Despite its mathematical simplicity, the LSA topics were less coherent compared to LDA, likely due to its lack of probabilistic modeling.

Table 6. Top 8 topics extracted using Latent Semantic Analysis (LSA). Each row lists the top keywords per topic.

Topic	Top Keywords
Topic 1	case, ipad, cover, keyboard, stand, protection, fits
Topic 2	drive, keyboard, macbook, pro, mac, windows, laptop
Topic 3	tv, samsung, screen, hdmi, player, quality, remote
Topic 4	keyboard, keys, bluetooth, mouse, wireless, key, logitech
Topic 5	cable, tv, works, apple, ipad, usb, samsung, charge
Topic 6	tablet, galaxy, samsung, screen, tab, protector, note
Topic 7	case, keyboard, cable, galaxy, samsung, tablet, tab, fit
Topic 8	macbook, great, pro, product, apple, price, perfectly

#### 4.4. Non-negative Matrix Factorization (NMF)

NMF [11] factorizes the TF-IDF matrix into two lower-rank non-negative matrices. It results in additive topic representations, making the output interpretable and often more sparse. The topics extracted from NMF included:

Table 7. Top 9 topics extracted using Latent Semantic Analysis (NMF). Each row lists the top keywords per topic.

Topic	Top Keywords
Topic 1	Samsung, amazon, problem, support, screen
Topic 2	Ipad, Cover, Case, protection, Leather, Screen, laptop
Topic 3	Hard, Drive, SSD, external, windows, computer
Topic 4	Laptop, Lenovo, Thinkpad, battery, fits, life
Topic 5	Charger, charge, power, adaptor, iphone, port
Topic 6	tablet, galaxy, samsung, note, great, samsung tab
Topic 7	keyboar, logitech, mouse, wireless, bluetooth,
Topic 8	Cable, HDMI, Monitor, connect , adaptor
Topic 9	sound, quality, headphones, speaker , music, volume

#### 4.5. BERTopic: Transformer-Based Topic Modeling

BERTopic [7] combines transformer-based sentence embeddings (using BERT-like models) with HDBSCAN clustering and class-based TF-IDF for topic extraction. Among all tested models, BERTopic produced the most coherent and fine-grained topics.

Its ability to embed contextual semantics at the sentence level allowed it to group together nuanced feedback such as:

Table 8. Top 8 topics extracted using bertopic. Each row lists the top keywords per topic.

Topic	Top Keywords
Topic 1	ipad, case, cover, stand, it, the, is, leather
Topic 2	headphones, ear, these, sound, headset,
Topic 3	card, sandisk, sd, cards, reader, class,
Topic 4	charge, charger, charging, phone, usb, battery
Topic 5	mouse, scroll, wheel, logitech, mice,
Topic 7	keyboar, logitech, mouse, wireless, bluetooth,
Topic 8	camera, lens, zoom, canon, video, focus, picture

#### 4.6. Summary and Insights

#### 4.7. Summary and Insights

Overall, topic modeling highlighted common themes in user reviews that align with our sentiment classification results. While LDA and NMF provided interpretable results with minimal preprocessing, BERTopic stood out in identifying subtle, context-rich groupings.

### 5. Conclusion

In this study, we conducted a comprehensive sentiment analysis on technology product reviews, focusing primarily on brands such as Apple and Samsung. We experimented

with a diverse set of modeling techniques, ranging from classical machine learning baselines to deep learning architectures and modern transformer-based language models. Our pipeline also included an extensive topic modeling analysis to uncover the latent thematic structure in user opinions.

Through rigorous experimentation, we observed that traditional approaches such as Logistic Regression or Decision Trees struggled to model sentiment intricacies, particularly in neutral and context-dependent reviews. Shallow recurrent models like single-layer GRUs showed some promise but lacked the expressive power required for nuanced sentiment classification.

Advanced architectures such as BiLSTM with attention mechanisms, models enriched with GloVe embeddings, and focal loss-based training significantly improved performance. Among these, the best-performing deep learning model was the BiLSTM + Attention + GloVe + Focal Loss, achieving an accuracy of 70% and demonstrating robustness in handling imbalanced and noisy labels.

Transformer-based models such as DeBERTaV3-large and CardiffNLP-TweetEval offered the highest accuracy and F1-scores, showcasing the effectiveness of pre-trained contextual embeddings and large-scale language understanding. Additionally, topic modeling with LDA, NMF, and BERTopic helped reveal the prominent user concerns and brand-specific feedback patterns. From the topic modeling perspective, both BERTopic and NMF produced more coherent and generalizable topics than LDA, especially when capturing real-world brand and product discussions. Overall, our findings highlight the importance of dataset quality, modeling depth, and robust evaluation in building effective sentiment analysis pipelines applicable to practical domains such as product feedback mining, brand monitoring, and market analysis.

As a continuation of our previous work, we initially explored the same sentiment classification task on a custom dataset of 8,000 manually labeled Reddit comments related to the 2024 U.S. election. While that dataset yielded high accuracy (around 96%) using transformer models, it was relatively biased: clear and strong positive or negative comments dominated, while neutral examples were either scarce or generated artificially using large language models. This resulted in a simpler classification problem that did not reflect real-world complexity.

In contrast, our current work focused on a more general and challenging dataset composed of real user reviews, which better captured the ambiguity and variability of public sentiment. Last semester, we invested considerable effort in annotating this new dataset with a balanced distribution of sentiments, and the resulting model performance revealed deeper insights into the limitations of simpler models.

To support reproducibility and transparency, all the code implementations for each stage — including preprocessing, model training, evaluation, and topic modeling — have been publicly uploaded to a GitHub repository, complete with detailed documentation and executable examples.

Model	Validation Accuracy	Test Accuracy	F1-Score
DistilBERT	87.1%	91.6%	91.4%
RoBERTa-base	88.7%	90.4%	90.5%
RoBERTa-large	98.2%	97.5%	96.6%
DeBERTa-v3-small	88.8%	91.0%	91.0%
DeBERTa-large	82.0%	89.3%	87.0%

Table 9. Comparison of model performance across different architectures on the U.S. Election 2024 dataset.

In addition to transformer models, we also evaluated custom deep learning architectures on the U.S. Election 2024 dataset. A Bidirectional LSTM model with GloVe embeddings and an attention mechanism achieved a strong test accuracy of 94.3%, outperforming the GRU-based variant, which reached 88.0%. These results demonstrate that, when supported by quality embeddings and attention mechanisms, recurrent models can still offer competitive performance on clean and well-separated sentiment datasets.

Deep Learning Model	Test Accuracy
BiLSTM + GloVe + Attention	94.3%
GRU + GloVe + Attention	88.0%

Table 10. Performance of deep learning models on the U.S. Election 2024 dataset.

## Work Contribution and Time Breakdown

### Team Members: Mohammad and Bahar

Below is a breakdown of the main activities and the time spent by each team member. Total effort is approximately **34 hours for Mohammad** and **36 hours for Bahar**, in line with course guidelines.

Task	Mohammad (hrs)	Bahar (hrs)
<b>Dataset Collection and Labeling</b>		
Reddit dataset (USA Election) creation + labeling	8	8
Amazon reviews filtering and new dataset prep	2	3
Manual labeling + synthetic neutrals	2	2
<b>Sentiment Classification</b>		
Transformer models (DeBERTa, RoBERTa, XLNet, CardiffNLP)	5	5
Deep learning models (LSTM, GRU, BiLSTM + Attention + GloVe)	8	4
FastText implementation + evaluation	1	2
Classical ML models (LogReg, Decision Tree, RF)	2	4
<b>Topic Modeling</b>		
LDA, LSA, NMF implementations + analysis	3	3
BERTopic analysis + figures	2	3
Topic interpretation and LaTeX tables	1	1
<b>Report Writing and Submission</b>		
LaTeX formatting, tables, figures	2	2
Writing (intro, methods, results, conclusion)	3	3
Final edits, polishing, GitHub prep	1	0
<b>Total Hours</b>	<b>34</b>	<b>36</b>

**Infrastructure and Code:** For the custom dataset (U.S. Election), we performed all training and fine-tuning on Google Colab using GPUs. For the large-scale Apple-Samsung dataset (approx. 70k reviews), we utilized the university compute cluster. All models, log files, and evaluation outputs are documented and publicly accessible in the accompanying GitHub repository.

**Note:** *Amir Sadeghi* was another teammate who contributed to the labeling of the USA Election dataset. He completed this course in the previous semester and agreed to help with initial annotation efforts. Approximately **9 hours** of the Reddit dataset labeling was done by him.

## Appendix: Mathematical Foundations

### Recurrent Neural Networks

#### LSTM: Long Short-Term Memory

LSTMs introduce gating mechanisms to control information flow through time. The main equations governing an LSTM cell are:

$$f_t = \sigma(W_f[h_{t-1}, x_t] + b_f) \quad (\text{Forget Gate}) \quad (1)$$

$$i_t = \sigma(W_i[h_{t-1}, x_t] + b_i) \quad (\text{Input Gate}) \quad (2)$$

$$\tilde{C}_t = \tanh(W_C[h_{t-1}, x_t] + b_C) \quad (\text{Candidate Memory}) \quad (3)$$

$$C_t = f_t \cdot C_{t-1} + i_t \cdot \tilde{C}_t \quad (\text{Cell State Update}) \quad (4)$$

$$o_t = \sigma(W_o[h_{t-1}, x_t] + b_o) \quad (\text{Output Gate}) \quad (5)$$

$$h_t = o_t \cdot \tanh(C_t) \quad (\text{Hidden State}) \quad (6)$$

## GRU: Gated Recurrent Unit

GRUs simplify LSTMs by using only reset and update gates:

$$z_t = \sigma(W_z[h_{t-1}, x_t] + b_z) \quad (\text{Update Gate}) \quad (7)$$

$$r_t = \sigma(W_r[h_{t-1}, x_t] + b_r) \quad (\text{Reset Gate}) \quad (8)$$

$$\tilde{h}_t = \tanh(W_h[r_t \cdot h_{t-1}, x_t] + b_h) \quad (9)$$

$$h_t = (1 - z_t) \cdot h_{t-1} + z_t \cdot \tilde{h}_t \quad (10)$$

## Transformer Models

### Self-Attention Mechanism

Transformers replace recurrence with self-attention. Given query  $Q$ , key  $K$ , and value  $V$  matrices, attention is computed as:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}}\right)V \quad (11)$$

Where:

- $Q, K, V$  are learned linear projections of input tokens.
- $d_k$  is the dimension of the key vectors (used for scaling).

### BERT and Variants

**BERT** is pre-trained using:

- **Masked Language Modeling (MLM)**: Predicts randomly masked tokens.
- **Next Sentence Prediction (NSP)**: Determines sentence continuity.

**RoBERTa** removes NSP and trains longer with more data. **DistilBERT** is a smaller, faster version of BERT via knowledge distillation.

### Topic Modeling Algorithms

#### Latent Dirichlet Allocation (LDA)

LDA models documents as mixtures of topics, where each topic is a distribution over words.

#### Generative Process:

1. For each document  $d$ :
  - (a) Draw topic proportions  $\theta_d \sim \text{Dirichlet}(\alpha)$ .
  - (b) For each word:
    - i. Draw a topic  $z \sim \text{Multinomial}(\theta_d)$ .
    - ii. Draw word  $w \sim \text{Multinomial}(\phi_z)$ .

#### Joint Probability Objective:

$$P(W, Z, \Theta, \Phi | \alpha, \beta) = \prod_{d=1}^D P(\theta_d | \alpha) \prod_{n=1}^{N_d} P(z_{d,n} | \theta_d) P(w_{d,n} | \phi_{z_{d,n}})$$

Posterior inference is approximated using variational inference or Gibbs sampling.

## Latent Semantic Analysis (LSA)

LSA uses Singular Value Decomposition (SVD) on the term-document matrix  $A \in \mathbb{R}^{m \times n}$ :

$$A = U \Sigma V^\top$$

Where:

- $U$ : term-topic matrix
- $\Sigma$ : singular values
- $V$ : document-topic matrix

To reduce noise, only top  $k$  singular values are kept:

$$A_k = U_k \Sigma_k V_k^\top$$

## Non-Negative Matrix Factorization (NMF)

NMF factorizes a non-negative matrix  $A$  into:

$$A \approx WH$$

Where:

- $W \in \mathbb{R}^{m \times k}$ : term-topic matrix
- $H \in \mathbb{R}^{k \times n}$ : topic-document matrix

## Objective Function:

$$\min_{W, H \geq 0} \|A - WH\|_F^2$$

NMF provides additive and interpretable topic representations.

## References

- [1] Dzmitry Bahdanau, Kyunghyun Cho, and Yoshua Bengio. Neural machine translation by jointly learning to align and translate. *arXiv preprint arXiv:1409.0473*, 2014.
- [2] Francesco Barbieri, Luis Espinosa Anke, and Jose Camacho-Collados. Tweeteval: Unified benchmark and comparative evaluation for tweet classification. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: Findings*, pages 1644–1650, 2020.
- [3] David M Blei, Andrew Y Ng, and Michael I Jordan. Latent dirichlet allocation. *Journal of Machine Learning Research*, 3(Jan):993–1022, 2003.
- [4] Kyunghyun Cho, Bart van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. Learning phrase representations using rnn encoder–decoder for statistical machine translation. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 1724–1734, 2014.
- [5] Scott Deerwester, Susan T Dumais, George W Furnas, Thomas K Landauer, and Richard Harshman. Indexing by latent semantic analysis. *Journal of the American Society for Information Science*, 41(6):391–407, 1990.
- [6] Alex Graves and Jürgen Schmidhuber. Framewise phoneme classification with bidirectional lstm and other neural network architectures. *Neural Networks*, 18(5-6):602–610, 2005.
- [7] Maarten Grootendorst. Bertopic: Neural topic modeling with class-based tf-idf. <https://github.com/MaartenGr/BERTopic>, 2022.
- [8] Pengcheng He, Xiaodong Liu, Jianfeng Gao, and Weizhu Chen. Deberta: Decoding-enhanced bert with disentangled attention. *arXiv preprint arXiv:2006.03654*, 2021.
- [9] Sepp Hochreiter and Jürgen Schmidhuber. Long short-term memory. *Neural computation*, 9(8):1735–1780, 1997.
- [10] Armand Joulin, Edouard Grave, Piotr Bojanowski, and Tomas Mikolov. Bag of tricks for efficient text classification. In *Proceedings of the 15th Conference of the European Chapter of the Association for Computational Linguistics*, pages 427–431, 2017.



- [11] Daniel D Lee and H Sebastian Seung. Learning the parts of objects by non-negative matrix factorization. *Nature*, 401(6755):788–791, 1999.
- [12] Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. Roberta: A robustly optimized bert pretraining approach. *arXiv preprint arXiv:1907.11692*, 2019.
- [13] Zhilin Yang, Zihang Dai, Yiming Yang, Jaime Carbonell, Ruslan Salakhutdinov, and Quoc V Le. Xlnet: Generalized autoregressive pretraining for language understanding. In *Advances in Neural Information Processing Systems*, volume 32, 2019.